

Contents

Assessment	1
The principle behind all three	1
Academic integrity	1

Assessment

Three components, each testing something the others cannot.

Component	What it tests	When
Weekly exercises	That you can apply each method soon after meeting it	Set every Friday, due the next
Final project	That you can run an honest end-to-end study on an unfamiliar dataset	Brief at Lesson 1, submitted after Lesson 10
Written exam	That you understand <i>why</i> the methods work, not only how to call them	End of course

Weights are confirmed with the degree programme; see [../Course/SYLLABUS.md](#).

The principle behind all three

A student who can program will produce a model that appears to work within weeks. The difficult, valuable skill is telling that apart from a model that actually generalises.

So all three components weight **methodology** heavily. A project with a modest accuracy and a scrupulous evaluation scores well. A project with a spectacular accuracy obtained by leaking the test set into feature engineering does not — that is the single most common way a machine learning result turns out to be worthless, and detecting it is a course learning outcome.

Academic integrity

Discussing ideas with other students is encouraged; the peer review step depends on it. Submitting work you did not write is not.

You may use libraries freely unless an exercise explicitly asks for a from-scratch implementation. You may use AI assistants to help you understand material or debug code — but you must be able to explain every line you submit, and the exam is closed book. A submission you cannot explain is treated as one you did not write.

Cite any source you drew on: datasets, papers, blog posts, code. A short “Sources” section at the end of a notebook is enough.