

Contents

Assessment	1
The principle behind it	1
Academic integrity	1

Assessment

One component, and it is the ten weekly exercises.

What	What it tests	When
Weekly exercises	That you can apply each method soon after meeting it, and can say why you did what you did	Set at the end of each lesson, discussed at the start of the next
Exam	The same thing, out loud, about one of those exercises drawn at random	End of course

Nothing is collected and nothing is marked week by week. At the exam **one of the ten is drawn** and you talk through your own notebook for it, so all ten are worth keeping and none is optional.

The principle behind it

A student who can program will produce a model that appears to work within weeks. The difficult, valuable skill is telling that apart from a model that actually generalises.

So the assessment weights **methodology** heavily. A modest accuracy with a scrupulous evaluation does well. A spectacular accuracy obtained by leaking the test set into feature engineering does not — that is the single most common way a machine learning result turns out to be worthless, and detecting it is a course learning outcome.

This is also why the assessment is a conversation rather than a submission. A notebook can be produced by someone else, or by a tool; an account of why each decision in it was made cannot.

Academic integrity

Discussing ideas with other students is encouraged. Submitting work you did not write is not.

You may use libraries freely unless an exercise explicitly asks for a from-scratch implementation. You may use AI assistants to help you understand material or debug code — but you must be able to explain every line you write. Work you cannot explain is treated as work you did not write — and since the exam draws one of the ten exercises for you to discuss, that is not a threat but a description of how the exercises are assessed.

Cite any source you drew on: datasets, papers, blog posts, code. A short “Sources” section at the end of a notebook is enough.