

Contents

Exam	1
The written paper	1
The discussion	1
What you are expected to be able to derive	2
Materials, and when they appear	2
How to prepare	3

Exam

Two stages on the same day: a **written paper**, closed book, two hours; then a **discussion** of one of the ten weekly exercises, drawn at random.

The exam tests what the project cannot: whether you understand *why* the methods work. The paper draws on the **handouts**, including the derivations. If you have only run the notebooks, you will not pass.

The written paper

Part	Questions	Marks	What it tests
A — Foundations	8 short answers	24	Definitions, statements, when a method applies
B — Derivations	2 of 3 offered	30	Reproduce and reason about a derivation from the handouts
C — Method selection	2 scenarios	26	Choose an approach for a described problem and justify it
D — Diagnosis	1 case	20	Given results, say what went wrong and how to check

Part D always presents a flawed study. Finding the flaw is a course learning outcome, so it is never worth fewer than a fifth of the marks.

The discussion

One of the ten weekly exercises is drawn, and you talk through **your own notebook** for it. **Bring all ten**, on a laptop or printed; an exercise you cannot produce is an exercise you did not do.

It is a conversation about work you have already done, not a second examination of the theory. Expect to be asked:

- what you actually measured, and why that quantity rather than another;
- which decision in the notebook you were least sure of, and what you did about it;
- what result would have told you the whole thing was wrong.

There are no marks for the score you obtained. There are marks for knowing why you obtained it, which is the same standard the exercises are written to and the same one the project is marked against.

Its weight beside the written paper is confirmed with the degree programme; see [../../Course/SYLLABUS.md](#).

What you are expected to be able to derive

From the handouts, unaided:

Lesson	Derivation
3	Normal equation from the least-squares objective; the gradient descent update; Ridge as constrained optimisation
4	The logistic model's cross-entropy from maximum likelihood; why squared error is a poor choice for classification
5	The bias-variance decomposition of expected prediction error
6	The SVM margin as an optimisation problem; what the kernel trick avoids computing
7	Entropy and information gain for a split; boosting as gradient descent in function space
8	PCA as an eigenvalue problem; the k-means objective and why the algorithm converges
9	Backpropagation for a two-layer network; why non-linear activations are necessary

You are not asked to memorise library APIs. You may be asked to write a few lines of pseudocode.

Materials, and when they appear

Neither of these exists yet. They are listed here so you know what is coming and do not go looking for it early; both will appear in this folder, and reach you with a `git pull` like everything else.

Material	What it is	When
<code>question_bank.md</code>	Questions by topic, in the shape of Parts A, C and D	Built up as the course runs, from Lesson 4 onwards
<code>sample_paper_1.md</code> , <code>sample_paper_2.md</code>	Two full papers with worked solutions	After Lesson 8, Friday 13 November 2026

Until they do, the quizzes are the closest thing there is, and they are already written: one per lesson, in `Lessons/NN_topic/Quizzes/`. See *How to prepare* below.

How to prepare

Keep the ten exercise notebooks somewhere you can find them, and re-read your own reasoning before the day rather than rewriting the code — the discussion is about the decisions, and those live in the markdown cells.

For the paper: work through the derivations with a pen rather than reading them. Then take the quiz for each lesson — they deliberately include questions that require reasoning about a derivation rather than recalling a fact, and until the sample papers appear in November they are the closest thing to Part B there is.