



Lesson 2 — Data: Exploration and Preparation

Technologies for Artificial Intelligence

Fabio Antonini — Università degli Studi dell'Aquila

Before we start

Exercise 1 was due at the start of today's lesson.

- The wine-quality workflow, end to end
- Marked on **methodology**, not on accuracy
- We will look at a few submissions' framing choices in a moment

What today builds on

- Lesson 1's rule: **nothing is learned before the split**
- Lesson 1's tool: `Pipeline`, so the rule is structural, not just discipline
- Today: apply both to data that is actually messy

Today

- Exploratory analysis
- Missing values
- Outliers
- Scaling and encoding
- Feature engineering and pipelines
- Leakage, in a form nothing warns you about

One dataset, all lesson long

2000 synthetic telecom customers. Will they **churn?**

- Numeric: tenure, monthly charges, age, support calls
- Categorical: contract type, region, zip code (493 levels)
- Missing values, outliers, and one column built to carry **no signal**

Look before you touch anything


`.info()`, `.describe()`, class balance, first — always.

2000 rows, 8 columns, mixed types


Churn rate: **19.4%** → baseline accuracy 80.6%

tenure ... Note: this is the slide.

Raw distributions, before any cleaning



Where the gaps are



Three reasons a value is missing

Three reasons a value can be missing

MCAR

P(missing)



depends on nothing

MAR

P(missing)



depends on an
OBSERVED column

MNAR

P(missing)



depends on the
MISSING value itself


age: MCAR

Some customers just left it blank. Nothing systematic.

- Safe to impute with a fixed statistic
- The bias question is not “is this fair” — it is “what does filling it cost”

What mean imputation actually costs

Every correlation involving the imputed column shrinks — even under MCAR.



num_support_calls: MAR

Missing more often for long-tenure customers — not by chance.

- Treating it as MCAR reweights the sample the wrong way
- A **missingness indicator** column preserves the signal even after filling

What to do about a gap


Drop rows — only if few, and only if MCAR

Drop the column — if it is missing too often to help

Impute — mean/median, most-frequent,

KNNImputer

Add a missing indicator



is MAR signal


Outliers, two rules

Two standard rules for flagging an unusual value — the z-score, and Tukey's fence built on the interquartile range (IQR):

$$z_i = \frac{x_i - \bar{x}}{s}$$



$$[Q_1 - 1.5 \text{ IQR}, \quad Q_3 + 1.5 \text{ IQR}]$$

Calibrated to agree — on a normal column



The two rules disagree on real data

What the IQR rule catches



What neither rule can tell you

`tenure_months` of **-3** is impossible. Neither rule flags it.

- -3 is not *extreme* relative to a column spanning 0-72
- It is not *unusual*. It is *invalid* — a different question entirely

Once something is flagged

Detection finds candidates. What to do next is a **domain** decision.

- **Cap** it at the fence (Winsorise)
- **Remove** the row
- **Correct** it, if the true value is recoverable
- **Leave** it — unusual is not the same as wrong

Notebook 1, live

Exploration, missingness mechanisms, both outlier rules — from scratch.

Break

Why scale?

tenure_months: 0-70. monthly_charges: 15-150.


Gradient descent takes **one step size, along every axis, every iteration.**

The Hessian is the feature covariance


For uncorrelated, centred features, H is diagonal — the variances sit on it.

The Hessian is the shape of the bowl

Unscaled: a narrow valley





Scaled: a round bowl



A 100:1 variance ratio, and what it costs




A 100:1 variance ratio, and what it costs



StandardScaler vs MinMaxScaler

Two standard choices, both fitted on training data only.

tenure_months (x) against monthly_charges (y)



The dummy variable trap

One-hot encode all k categories AND fit an intercept:

$$\sum_{j=1}^k \text{dummy}_j = 1 = \text{intercept}$$

Rank deficient by exactly one

Why all k dummies plus an intercept cannot work

contract_type	m2m	1y	2y	sum	intercept
month-to-month	1	0	0	= 1	1
one-year	0	1	0	= 1	1
two-year	0	0	1	= 1	1

the k dummy columns always sum to the intercept column — rank deficient by exactly one

Ordinal vs target encoding




Ordinal

- Categories $\rightarrow$ integers
- Only valid if a real **order** exists
- `low < medium < high`

Target

- Category $\rightarrow$ a statistic of y
- Compresses any cardinality to one column
- The dangerous one — next

Three encodings, three claims




The curse of dimensionality, briefly

One-hot `zip_code`: **493** new columns, mostly zero.

- Distance-based methods (Lesson 6): everything looks equidistant
- More parameters to estimate, same sample size
→ more variance

zip_code: 493 levels




Restating Lesson 1's argument

f includes **every** learned parameter — not just “the model part”.

$$\mathbb{E}_{T \sim \mathcal{D}^m} [\hat{R}_T(f)] = R(f) \quad \text{only if } f \text{ is independent of } T$$

ColumnTransformer + Pipeline



The code

```
prep = ColumnTransformer([  
    ("num", num_pipe, num_cols),  
    ("cat", cat_pipe, cat_cols),  
])  
model = Pipeline([("prep", prep), ("clf",  
clf)])  
model.fit(X_train, y_train)
```


The model, on churn

Baseline: **0.806**

Model (numeric + low-card categorical, no zip):
0.820 accuracy, **0.752** AUC — area under the
receiver operating characteristic curve

Modest accuracy gain. Why?

Confusion matrix (test set)			
True label	stayed	churned	
	398	5	
stayed			
churned	85	12	
		Predicted label	
		stayed	churned
		85	12

Feature engineering

Ratios, interactions, binning — a linear model cannot invent these itself. AUC: 0.752 → 0.754, a modest but real gain.

$$\text{charge_per_tenure} = \frac{\text{monthly_charges}}{\text{tenure_months} + 1}$$

Notebook 2, live

Scaling from scratch, the dummy trap,
ColumnTransformer, Pipeline.

Same rule, broken three ways

Same rule, broken three ways — only the first one looks like a bug

Lesson 1

select features on the full dataset

visible — an obvious extra step

Today, leak 1

impute a missing value on the full dataset

invisible — looks like ordinary cleaning

Today, leak 2

encode a category on the full dataset


invisible — two unremarkable lines of code

Leak 1 — impute before splitting

`KNNImputer` fills a gap using **other rows' values**.
Fit it on train + test, and some of those rows are in the test set.

The smoking gun


98 of 128 training rows with missing age had a **test-set row** among their five nearest neighbours. Nothing raised an error.



Leak 2 — encode before splitting

Replace `zip_code` with the **average churn rate of its own customers**, computed before anybody has split anything.

A column with no real signal, encoded three ways




Why small groups make it worse

The “leave-in” encoding differs from the honest leave-one-out encoding by:

$$\bar{y}_c - \bar{y}_c^{(-i)} = \frac{y_i - \bar{y}_c^{(-i)}}{n_c}$$

The leak shrinks as $1/n_c$



The fix

`sklearn.preprocessing.TargetEncoder` — cross-fitted inside the pipeline.

Each row's encoded value comes from **other** folds, never its own label.

The rule, restated

Everything that **learns from data** belongs inside the fold it is fitted on.

A mean. A median. A set of neighbours. A per-category average. A set of bin edges.

Notebook 3, live

Trace the imputation leak. Watch the encoding leak manufacture 0.89 from noise.

What we did today

- Explored before transforming — types, gaps, shape, correlations
- Diagnosed *why* values are missing, not just *how many*
- Two outlier rules, and what neither can tell you
- Scaling, encoding, engineering — all inside Pipeline
- Two leaks that look nothing like leaking

Homework — due Friday 9 October

Exercises/

02_data_exploration_and_preparation.md

Build the full preparation pipeline yourself, and **find a leak on purpose** before fixing it.

Before next week

- Work the three notebooks in order
- Read the handout — the derivations are examinable
- Take the quiz
- **Do the exercise**

Next: regression — the first model we derive completely.