



Lesson 2: Data Exploration and Preparation

Technologies for Artificial Intelligence

Fabio Antonini, Università degli Studi dell'Aquila

Before we start

Exercise 1 was due at the start of today's lesson.

- The wine-quality workflow, end to end
- Marked on **methodology**, not on accuracy
- We will look at a few submissions' framing choices in a moment

What today builds on

- Lesson 1's rule: **nothing is learned before the split**
- Lesson 1's tool: `Pipeline`, so the rule is structural, not just discipline
- Today: apply both to data that is actually messy

Today

- Exploratory analysis
- Missing values
- Outliers
- Scaling and encoding
- Feature engineering and pipelines
- Leakage, in a form nothing warns you about

One dataset, all lesson long

2000 synthetic telecom customers. Will they **churn?**





- Numeric: tenure, monthly charges, age, support calls
- Categorical: contract type, region, zip code (493 levels)
- Missing values, outliers, and one column built to carry **no signal**

Look before you touch anything


- `.info()`, `.describe()`, class balance, first: always.
- 2000 rows, 8 columns, mixed types
- Churn rate: **19.4%** → baseline accuracy 80.6%
- `tenure_months` max: 999. Not a typo in the slide.

Two of these four have no shape left to read

Raw distributions, before any cleaning



Where the gaps are



Three reasons a value is missing

Three reasons a value can be missing

MCAR

P(missing)



depends on nothing

MAR

P(missing)



depends on an
OBSERVED column

MNAR

P(missing)



depends on the
MISSING value itself


age: MCAR

Some customers just left it blank. Nothing systematic.

- Safe to impute with a fixed statistic
- The bias question is not “is this fair”: it is “what does filling it cost”

What mean imputation actually costs

Every correlation involving the imputed column shrinks, even under MCAR.



num_support_calls: MAR

Missing more often for long-tenure customers, not by chance.


- Treating it as MCAR reweights the sample the wrong way
- A **missingness indicator** column preserves the signal even after filling

What to do about a gap

- **Drop rows:** only if few, and only if MCAR
- **Drop the column:** if it is missing too often to help
- **Impute:** mean/median, most-frequent, `KNNImputer`
- **Add a missingness indicator:** keeps MAR signal alive

Choosing among the four

Which one you pick is a question about why the data is missing



the mechanism decides, and every option except dropping the column learns a statistic, so it belongs inside the training fold



Rule 1: distance in standard deviations, flag beyond $k = 3$

$$Z_i = \frac{X_i - \bar{X}}{S}$$

Rule 2: distance measured in quartiles



$$[Q_1 - 1.5 \text{ IQR}, \quad Q_3 + 1.5 \text{ IQR}]$$

Same neighbourhood, not the same answer



The robust rule is the one that got it wrong

What the IQR rule catches



What neither rule can tell you

`tenure_months` of **-3** is impossible. Neither rule flags it.

- -3 is not *extreme* relative to a column spanning 0-72
- It is not *unusual*. It is *invalid*: a different question entirely

Once something is flagged

Detection finds candidates. What to do next is a **domain** decision.

- **Cap** it at the fence (Winsorise)
- **Remove** the row
- **Correct** it, if the true value is recoverable
- **Leave** it: unusual is not the same as wrong

Notebook 1, live

Exploration, missingness mechanisms, both outlier rules: from scratch.

Break

Why scale?


tenure_months: 0-72. monthly_charges: 15-128.

Gradient descent takes **one step size, along every axis, every iteration.**


A ravine, not a bowl

One stride length, two very different directions

Unscaled: a narrow valley





Scaled: a round bowl



A 100:1 variance ratio, and what it costs


A 100:1 variance ratio, and what it costs



StandardScaler vs MinMaxScaler

Two standard choices, both fitted on training data only.

tenure_months (x) against monthly_charges (y)



All k dummies plus an intercept
cannot both be free

$$\sum_{j=1}^k \text{dummy}_j = 1 = \text{intercept}$$

Rank deficient by exactly one

Why all k dummies plus an intercept cannot work

contract_type	m2m	1y	2y	sum	intercept
month-to-month	1	0	0	= 1	1
one-year	0	1	0	= 1	1
two-year	0	0	1	= 1	1

the k dummy columns always sum to the intercept column: rank deficient by exactly one

Ordinal vs target encoding




Ordinal

- Categories $\rightarrow$ integers
- Only valid if a real **order** exists
- `low < medium < high`

Target

- Category $\rightarrow$ a statistic of y
- Compresses any cardinality to one column
- The dangerous one comes next

Three encodings, three claims





The curse of dimensionality, briefly

One-hot `zip_code`: **493** new columns, mostly zero.

- Distance-based methods (Lesson 6): everything looks equidistant
- More parameters to estimate, same sample size
→ more variance

zip_code: 493 levels

zip_code, one-hot encoded: 493 columns for 2000 rows




Restating Lesson 1's argument

f includes **every** learned parameter, not just “the model part”.

$$\mathbb{E}_{T \sim D^m} [\hat{R}_T(f)] = R(f) \quad \text{only if } f \text{ is independent of } T$$

ColumnTransformer + Pipeline

Every learned parameter comes from inside this box



The code


```
prep = ColumnTransformer([
    ("num", num_pipe, num_cols),
    ("cat", cat_pipe, cat_cols),
])
model = Pipeline([
    ("prep", prep),
    ("clf", clf),
])
model.fit(X_train, y_train)
```

The model, on churn

- Baseline: **0.806**
- Model (numeric + the low-cardinality categoricals, no zip): **0.820** accuracy, **0.751** AUC (area under the receiver operating characteristic curve)

Modest accuracy gain. Why?

Where the errors fall



Feature engineering

Ratios, interactions, binning: a linear model cannot invent these itself. AUC: 0.751 $\rightarrow$ 0.754. **Two thousandths: one split cannot tell that from noise.**

$$\text{charge_per_tenure} = \frac{\text{monthly_charges}}{\max(\text{tenure_months}, 0) + 1}$$

Notebook 2, live

Scaling from scratch, the dummy trap,
ColumnTransformer, Pipeline.

Same rule, broken three ways

Same rule, broken three ways: only the first one looks like a bug

Lesson 1

select features on the full dataset

visible: an obvious extra step

Today, leak 1

impute a missing value on the full dataset

invisible: looks like ordinary cleaning

Today, leak 2

encode a category on the full dataset

invisible: two unremarkable lines of code

Leak 1: impute before splitting


`KNNImputer` fills a gap using **other rows' values**.
Fit it on train + test, and some of those rows are in the test set.

The smoking gun

94 of 128 training rows with missing age had a **test-set row** among the five donors used to fill it. Nothing raised an error.

Leakage, counted


94 of 128 imputed values touched the test set



Leak 2: encode before splitting

Replace `zip_code` with the **average churn rate of its own customers**, computed before anybody has split anything.


A column with no real signal, encoded three ways



Your own disagreement, divided by
your group size

$$\bar{y}_c - \bar{y}_c^{(-i)} = \frac{y_i - \bar{y}_c^{(-i)}}{n_c}$$

The leak shrinks as $1/n_c$



The fix

`sklearn.preprocessing.TargetEncoder`: **CROSS-**fitted inside the pipeline.

Each row's encoded value comes from **other** folds, never its own label.

The rule, restated

Everything that **learns from data** belongs inside the fold it is fitted on.

A mean. A median. A set of neighbours. A per-category average. A set of bin edges.

Notebook 3, live

Trace the imputation leak. Watch the encoding leak manufacture 0.89 from noise.

What we did today

- Explored before transforming: types, gaps, shape, correlations
- Diagnosed *why* values are missing, not just *how many*
- Two outlier rules, and what neither can tell you
- Scaling, encoding, engineering: all inside Pipeline
- Two leaks that look nothing like leaking

Homework: due Friday 9 October, 09:00

Exercises/

02_data_exploration_and_preparation.md

Build the full preparation pipeline yourself, and **find a leak on purpose** before fixing it.

Before next week

- Work the three notebooks in order
- Read the handout: the derivations are examinable
- Take the quiz
- **Do the exercise**

Next: regression, the first model we derive completely.