

UNIVERSITÀ
DEGLI STUDI
DELL'AQUILA

Lesson 4: Classification and Evaluation Metrics

Technologies for Artificial Intelligence

Fabio Antonini, Università degli Studi dell'Aquila

23 October 2026

Before we start

- Exercise 3, which we discuss now
- Ridge, Lasso, and the choice of the penalty strength

UNIVERSITÀ
DEGLI STUDI
DELL'AQUILA

Today: from a number to a decision

- Logistic regression: the model and why the cost had to change
- The confusion matrix, and why accuracy lies
- Precision, recall, F1, and the threshold nobody chose
- Receiver operating characteristic (ROC) curves, and where they mislead
- Choosing a threshold from money

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The question has changed

- Lesson 3: given a house, **how many euros?**
- Lesson 4: given a disk drive, **will it fail?**
- Answer with a probability: $0 \leq P(y = 1 \mid x) \leq 1$
- $y = 1$ if the drive failed, $y = 0$ if not; x holds its six counters

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The fleet

- **8,000 disk drives** in a data centre
- Six SMART counters each: Self-Monitoring, Analysis and Reporting Technology
- Label: did it fail within thirty days?
- `train_test_split(..., stratify=y)`: 6,000 to train, 2,000 to test

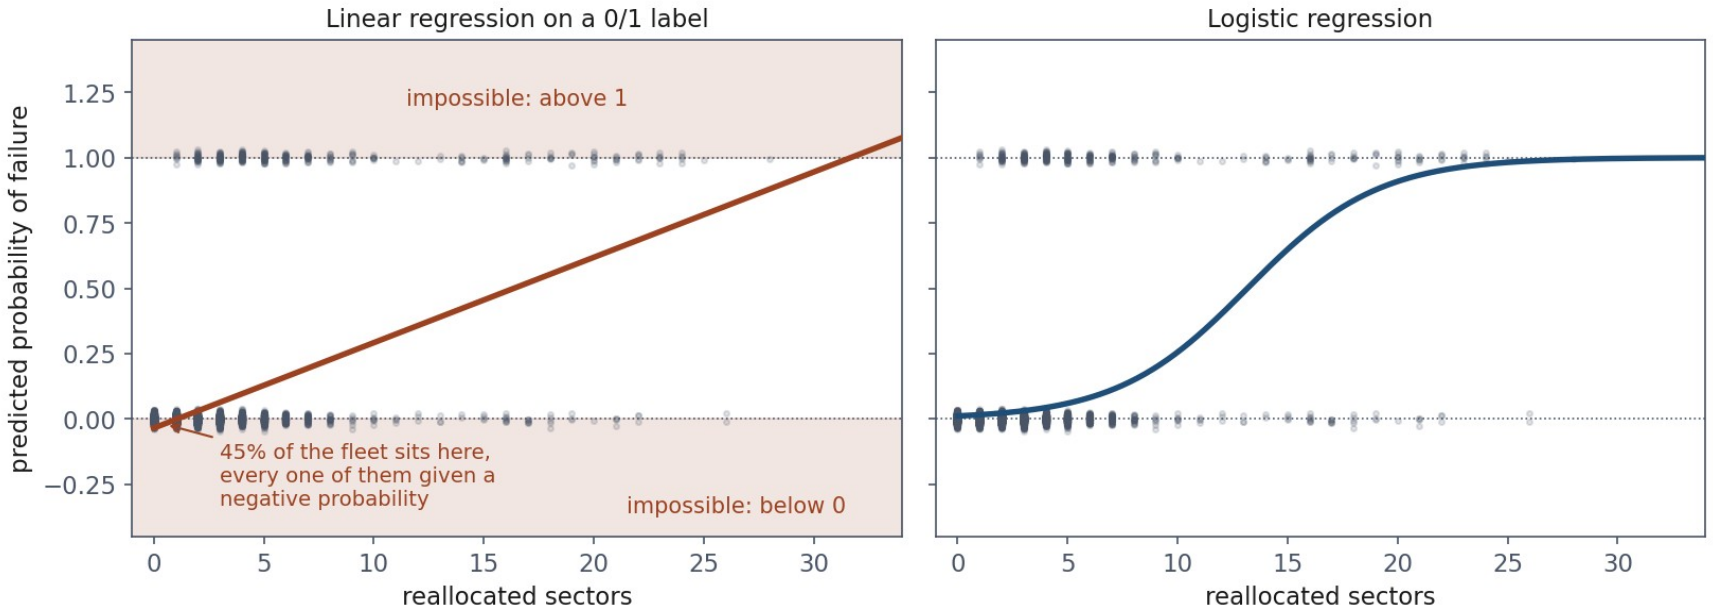
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

306

- **306 of the 8,000 drives failed.** 3.8% of the fleet
- Which means 96.2% did not

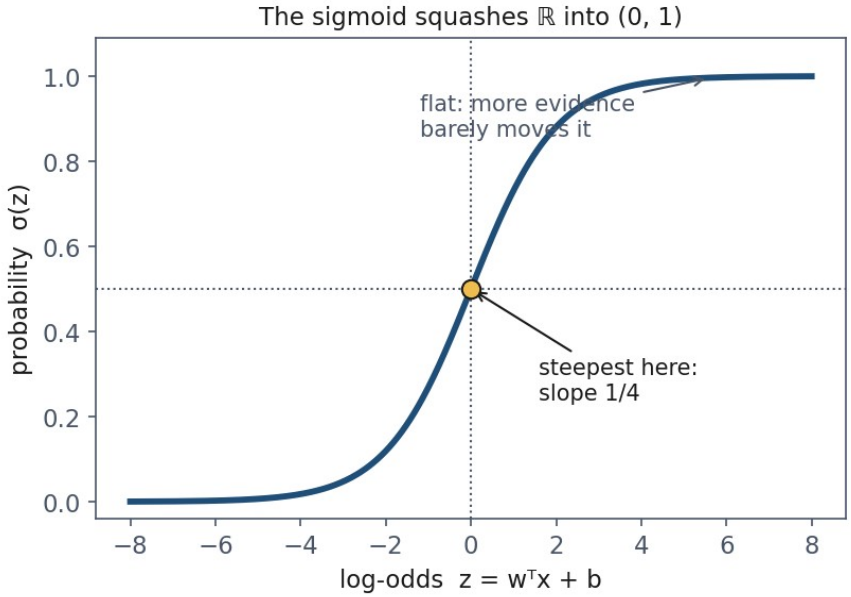
UNIVERSITÀ
DEGLI STUDI
DELL'AQUILA

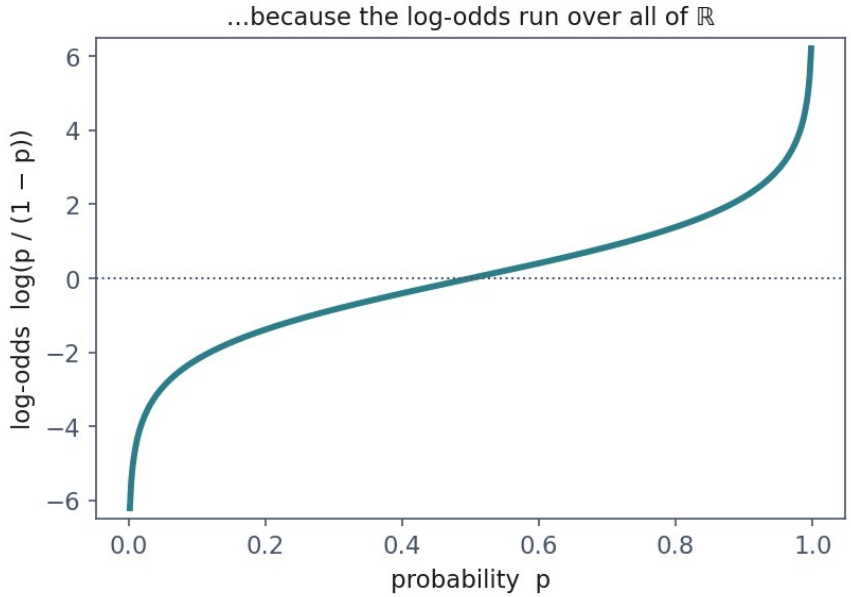
What goes wrong if we just fit a line



UNIVERSITÀ
DEGLI STUDI
DELL'AQUILA

The sigmoid





UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Why the log-odds

- A line gives any number; a probability, only [0, 1]
- Odds, $p \div (1 - p)$: no ceiling, from 0 to ∞
- Log of the odds: no floor either, $-\infty$ to $+\infty$
- $p = 0.01, 0.5, 0.99 \rightarrow$ log-odds $-4.60, 0, +4.60$
- Bonus: coefficients multiply the odds (slide 12)

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Log-odds of p = P(y = 1 | x): a line

$$\log \frac{p}{1-p} = w^\top x + b \quad \Longrightarrow \quad p = \frac{1}{1 + e^{-(w^\top x + b)}}$$

UNIVERSITÀ
DEGLI STUDI
DELL'AQUILA

The model, in one line

$$P(y = 1 \mid x) = \sigma(w^\top x + b), \qquad \sigma(z) = \frac{1}{1 + e^{-z}}$$

UNIVERSITÀ
DEGLI STUDI
DELL'AQUILA

What the coefficients mean

- The model is linear in the log-odds, log(p ÷ (1 − p)), so `exp(coefficient)` multiplies the **odds**, p ÷ (1 − p)
- Features standardised, so one unit = one standard deviation
- A coefficient of 1.80 means **six times the odds** of failing
- In code: `make_pipeline(StandardScaler(), LogisticRegression())`

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The truth, and what we recovered

Feature	True	Fitted	Odds ×
reallocated_sectors	1.80	1.66	5.28
spin_retry_count	1.15	1.17	3.22
read_error_rate	0.85	0.87	2.39
temperature_c	0.45	0.47	1.59
power_on_hours	0.35	0.38	1.46
seek_error_rate	**0.00**	**0.02**	**1.03**

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The intercept is the average drive’s chance

- Fitted intercept: **−6.09**
- σ(-6.09) = 0.0023: a 0.2% chance for a drive average on every counter
- The fleet fails at 3.8%: the failures come from the few drives far above average

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Fitting means minimising something, but what?

- Squared error worked for lesson 3
- Why not just apply it to the sigmoid’s output?

UNIVERSITÀ
DEGLI STUDI
DELL'AQUILA

Maximum likelihood, on three drives

Drive, what happened	Setting 1 of w, b: P(fail) → P(what happened)	Setting 2 of w, b: P(fail) → P(what happened)
A, failed	0.8 → **0.8**	0.3 → **0.3**
B, healthy	0.1 → **0.9**	0.6 → **0.4**
C, healthy	0.2 → **0.8**	0.5 → **0.5**
Multiplied: the likelihood	**0.576**	**0.060**
Maximum likelihood keeps	**setting 1**	

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Cross-entropy, also called log loss

$$L = -[\, y \log p + (1-y) \log(1-p)]$$

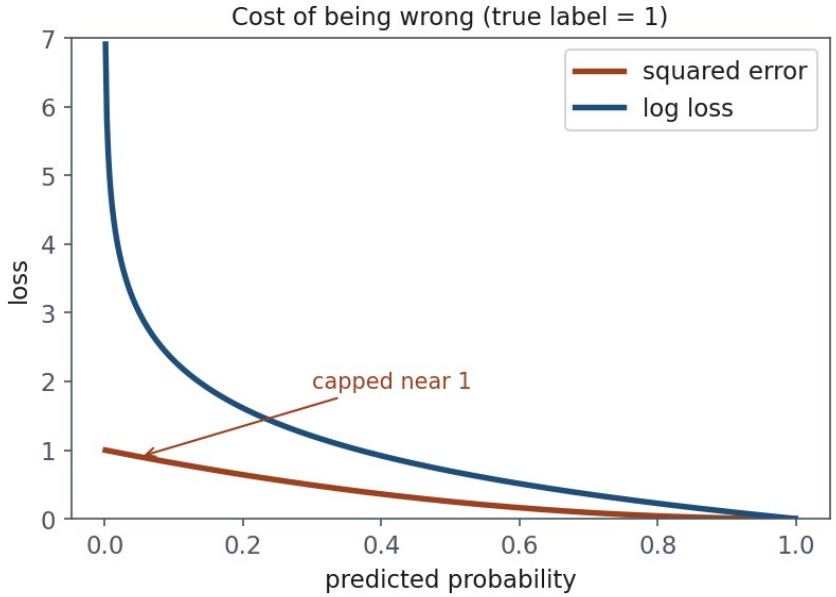
UNIVERSITÀ
DEGLI STUDI
DELL'AQUILA

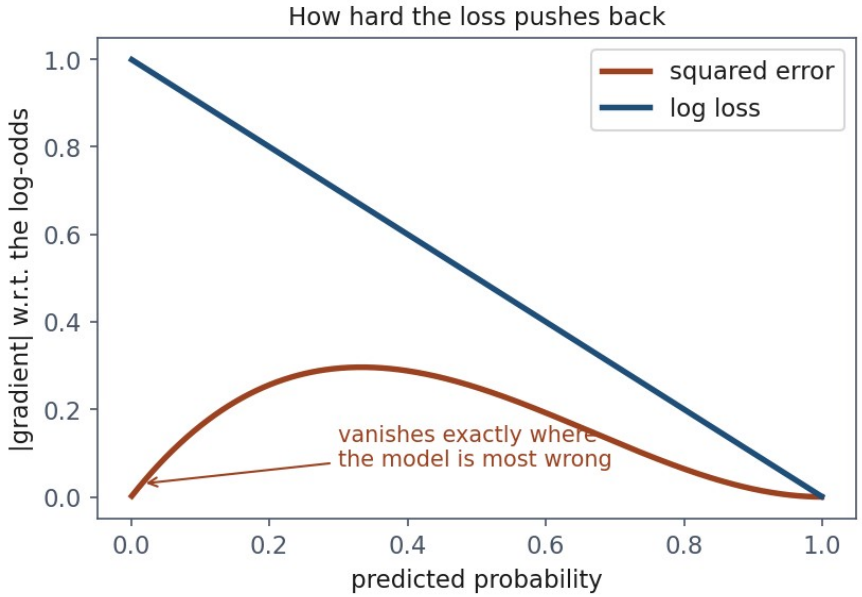
Lesson 3’s gradient, $\hat{y} = \sigma(w^T x + b)$

$$\frac{\partial J}{\partial w_j} = \frac{1}{m} \sum_{i=1}^{m} (\hat{y}^{(i)} - y^{(i)}) x_j^{(i)}$$

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

So why not squared error?





UNIVERSITÀ DEGLI STUDI DELL'AQUILA

And it is convex, which squared error is not

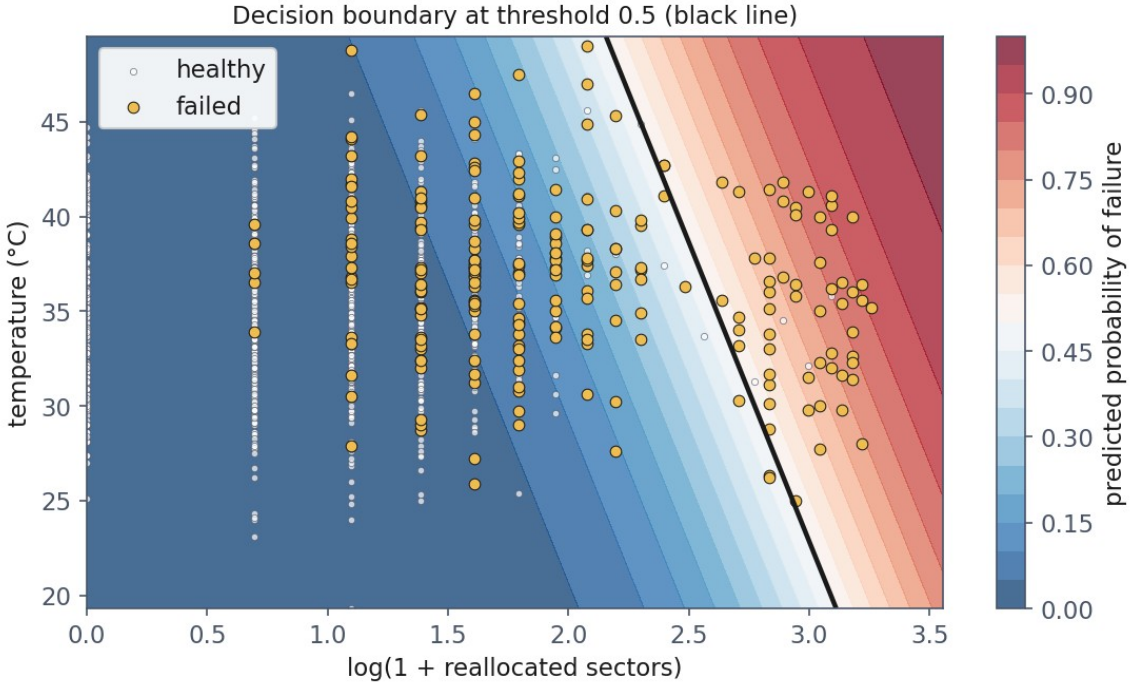
- Log loss is **convex** in the weights, so any local minimum is global
- Squared error through a sigmoid is not

UNIVERSITÀ
DEGLI STUDI
DELL'AQUILA

Notebook 1, live

section, and its slides	you call (→ the number to come back with)
a line on a 0/1 label, 7	`LinearRegression` → 3,606 negative probabilities
odds, and why not squared error, 8-20	NumPy only: the sigmoid, the two losses, their gradients
gradient descent, 18	`train_test_split`, `StandardScaler`, fifteen lines of NumPy → log loss 0.6931 to 0.0692
the decision boundary, 22	`make_pipeline(StandardScaler(), LogisticRegression())` on two features
the truth, recovered, 13-14	`LogisticRegression`, `.coef_`, `.intercept_` → seek_error_rate 0.02, intercept −6.09

The decision boundary



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Break

- Ten minutes

UNIVERSITÀ
DEGLI STUDI
DELL'AQUILA

Two models

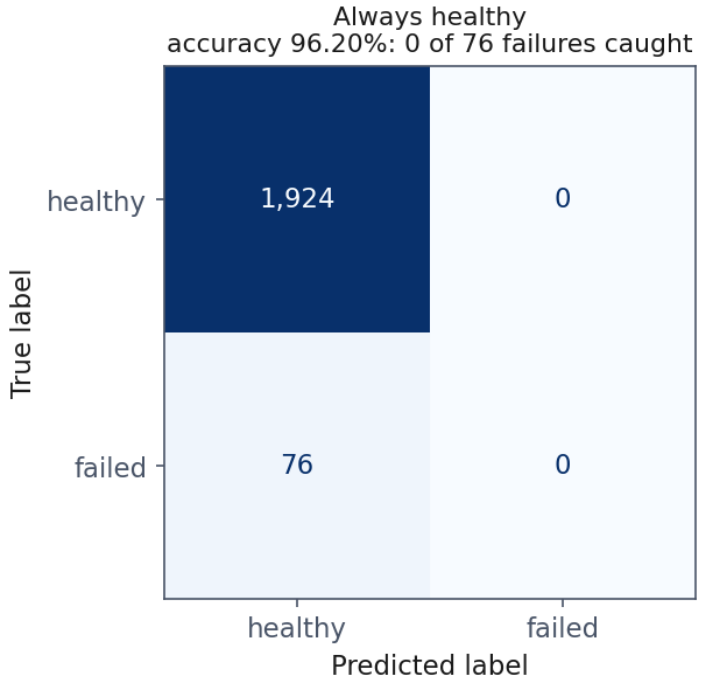
Model	Accuracy	Failures caught
Predict “healthy” for everything	**96.20%**	**0 of 76**
Logistic regression, threshold 0.5	97.70%	43 of 76

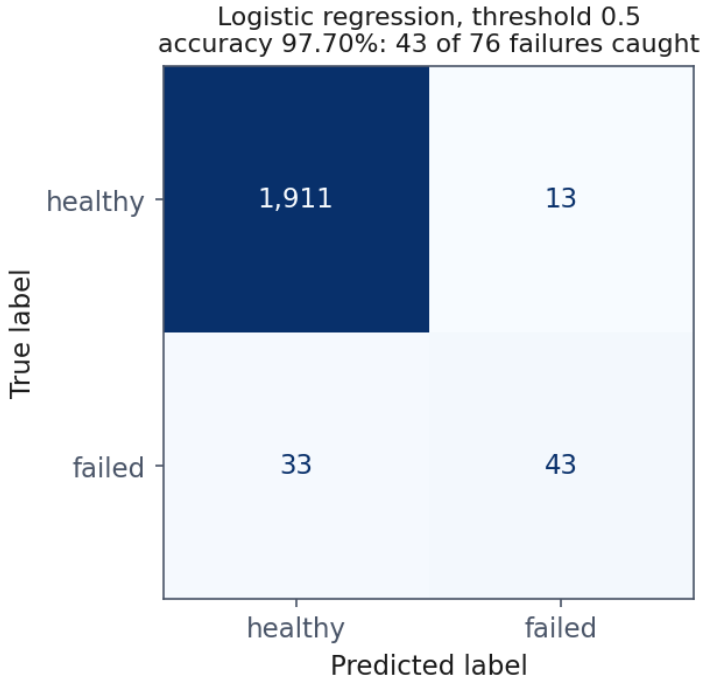
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Why everyone makes this mistake once

- Accuracy is the fraction of predictions that were right
- On a **balanced** problem that is exactly what you want
- The interesting class is almost always the **rare** one

Four numbers instead of one





UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The four outcomes, and what they cost

- **False positive**: we replace a healthy drive. A technician’s hour, €140
- **False negative**: a drive dies in service. Callout, rebuild, risk, €2,600

Precision and recall: look at the denominator

Precision = TP / (TP + FP)
"of the alarms we raised"

	predicted healthy	predicted failing
actually healthy	TN 1911	**FP** 13
actually failed	FN 33	**TP** 43

Recall = TP / (TP + FN)
"of the drives that failed"

	predicted healthy	predicted failing
actually healthy	TN 1911	FP 13
actually failed	**FN** 33	**TP** 43

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Our model, in those terms

- True and false positives and negatives: TP, FP, FN, TN
- Precision = TP ÷ (TP + FP) = 43 ÷ 56 = **0.768**: of the drives flagged
- Recall = TP ÷ (TP + FN) = 43 ÷ 76 = **0.566**: of the drives that failed
- Specificity = TN ÷ (TN + FP) = 1,911 ÷ 1,924 = 0.993
- In code: `precision_score`, `recall_score`

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

F1: why the harmonic mean

Model	Precision	Recall	Arithmetic mean	Harmonic, F1 = 2PR ÷ (P + R)
Flag every drive	0.038	1.000	0.519	**0.073**

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Read the minority class, not the average

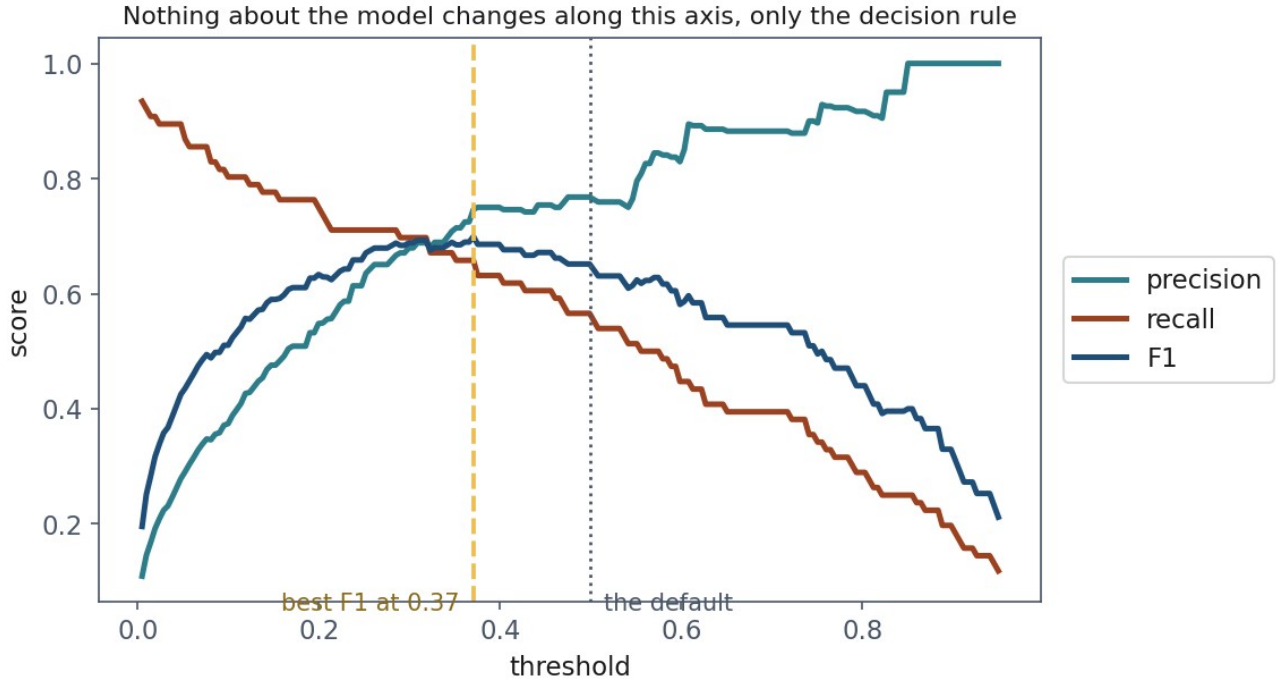
- `weighted avg` F1: **0.975** (dominated by the healthy class)
- `macro avg` F1: **0.820** (both classes get an equal say)
- In code: `classification_report(y_test, y_pred)` prints both rows

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Notebook 2, live

section, and its slides	you call (→ the number to come back with)
the number that should worry you, 24	`accuracy_score` → 96.20% for “always healthy”
the confusion matrix, 26	counts by hand, then `confusion_matrix`, `ConfusionMatrixDisplay` → TP 43
precision and recall, 28-31	`precision_score`, `recall_score`, `f1_score`, `classification_report`
the threshold sweep, 33-34	`predict_proba`, then the metrics at each threshold
the precision-recall curve, 35	`precision_recall_curve`, `average_precision_score` → 0.717 against 0.038

The threshold is a choice nobody made



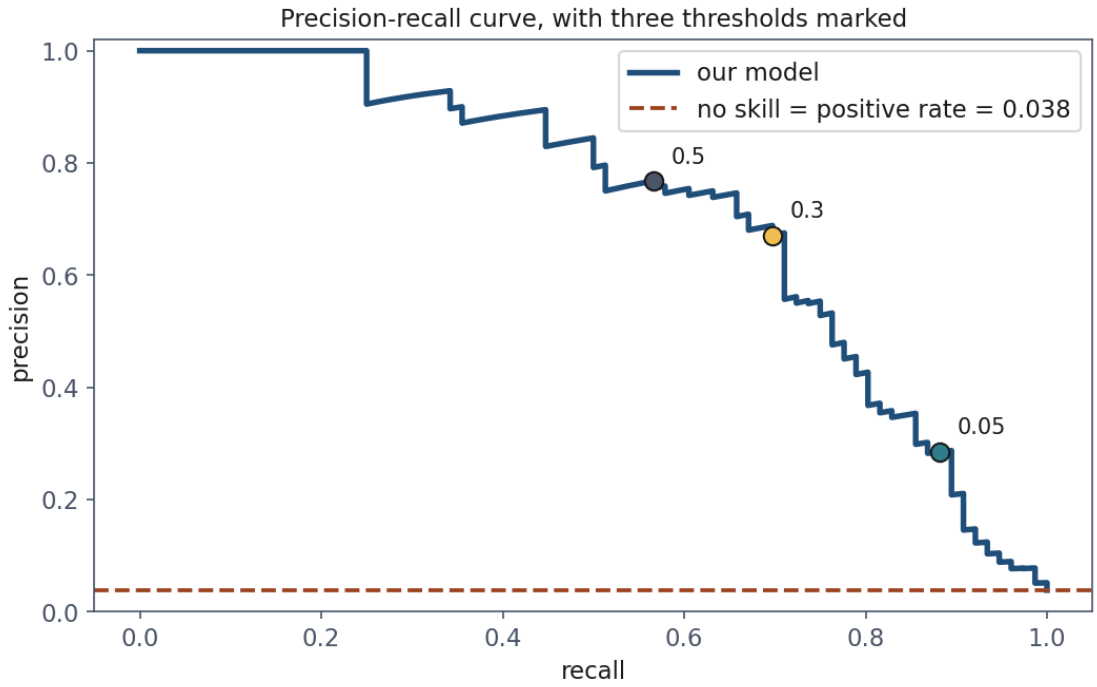
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

What the sweep actually costs

Threshold	TP	FP	FN	Precision	Recall	Accuracy
0.02	69	284	7	0.195	0.908	0.855
0.10	61	102	15	0.374	0.803	0.942
0.30	53	25	23	0.679	0.697	0.976
0.50	43	13	33	0.768	0.566	**0.977**
0.90	14	0	62	1.000	0.184	0.969

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The precision-recall curve



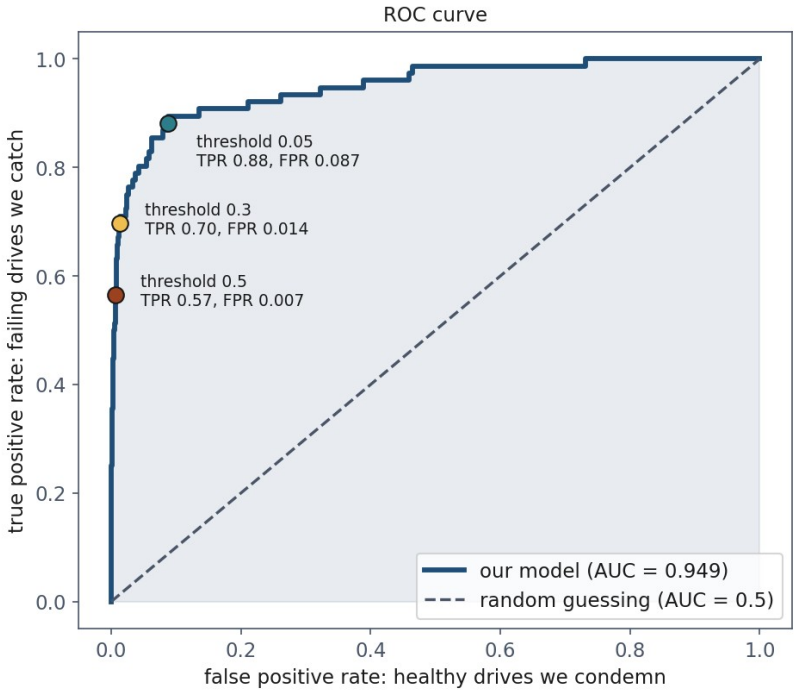
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Comparing models without choosing a threshold

- We keep evaluating at one threshold, but which?
- What if we score the whole **ranking** instead?

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The ROC curve



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

What the area means

- **Area under the curve (AUC) = the probability that a random failing drive scores above a random healthy one**
- Ours: 0.949
- In code: `roc_auc_score(y_test, scores)`

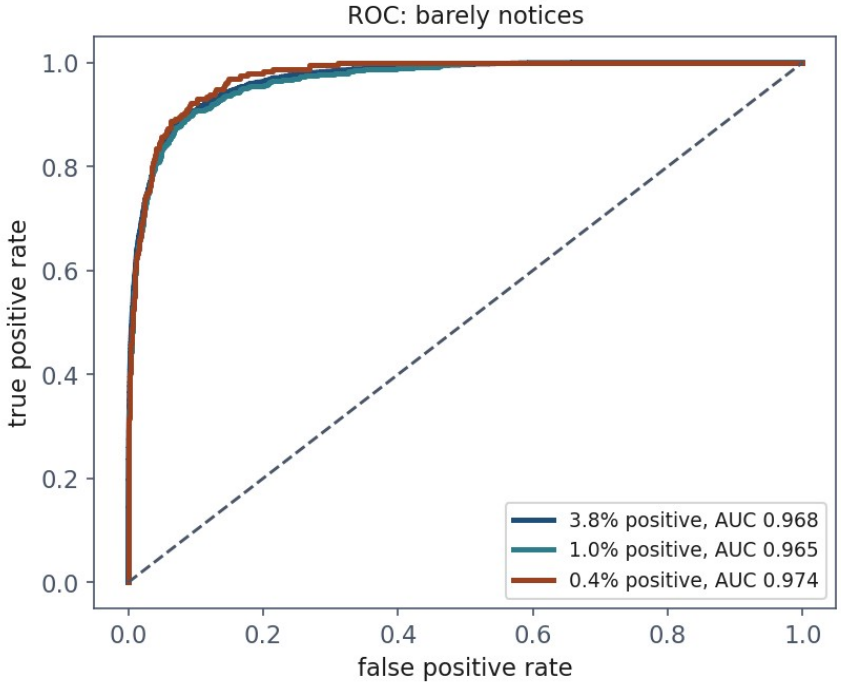
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

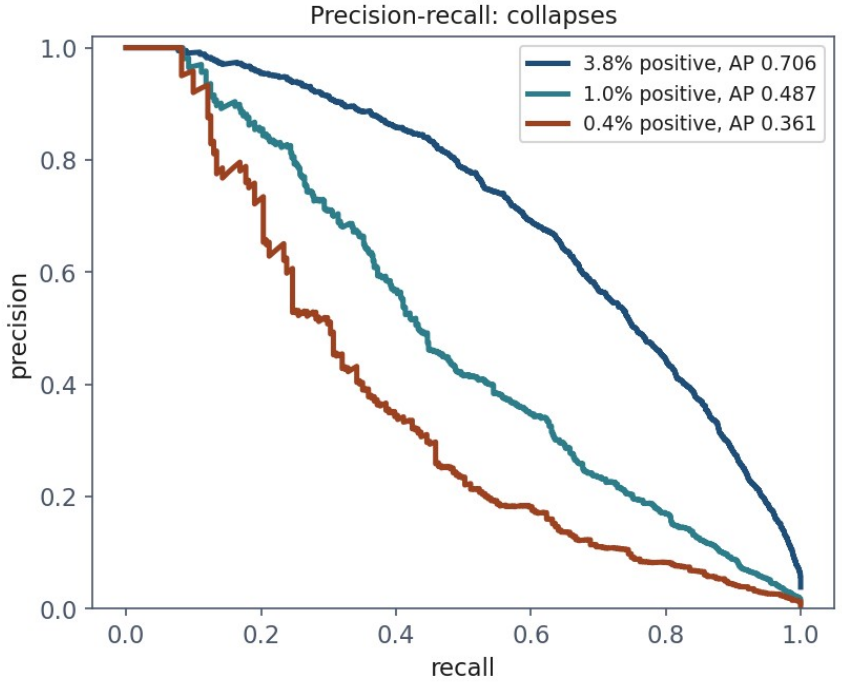
What AUC does not tell you

- Divide every predicted probability by 10
- AUC is **unchanged**; every decision changes

UNIVERSITÀ
DEGLI STUDI
DELL'AQUILA

Where ROC misleads





UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Why the two disagree

- False positive rate (FPR) divides by **all healthy drives**: an enormous, growing denominator
- Precision divides by **the alarms raised**, where the false ones dominate

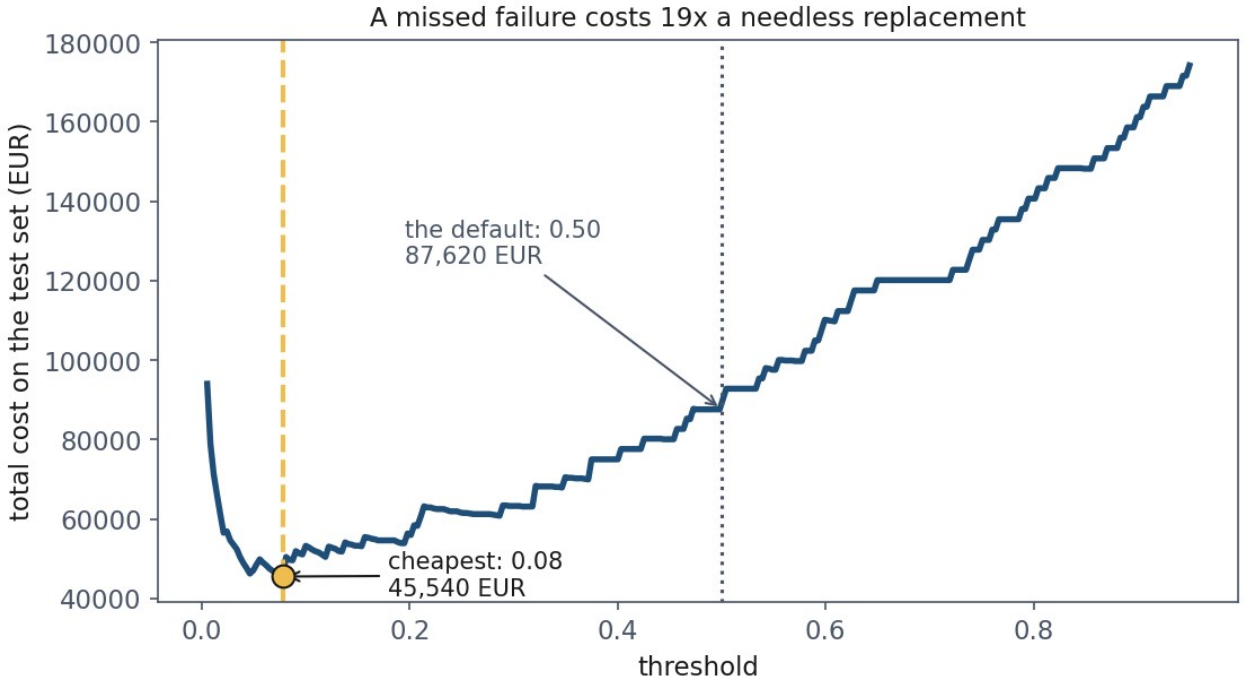
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The rule

- On an imbalanced problem, report the **precision-recall curve alongside AUC**
- Never instead of it, and never AUC alone
- Its area is the **average precision**, `average_precision_score`: 0.717 here, against a floor of 0.038

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Choosing the threshold from money



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

And it has a closed form

$$t^* = \frac{C_{FP}}{C_{FP} + C_{FN}} = \frac{140}{140 + 2600} = 0.051$$

UNIVERSITÀ
DEGLI STUDI
DELL'AQUILA

The caveat you must state

- We chose the threshold **by looking at the test set**
- That is precisely what lesson 1 forbade

Class weights: the same lever, pulled earlier

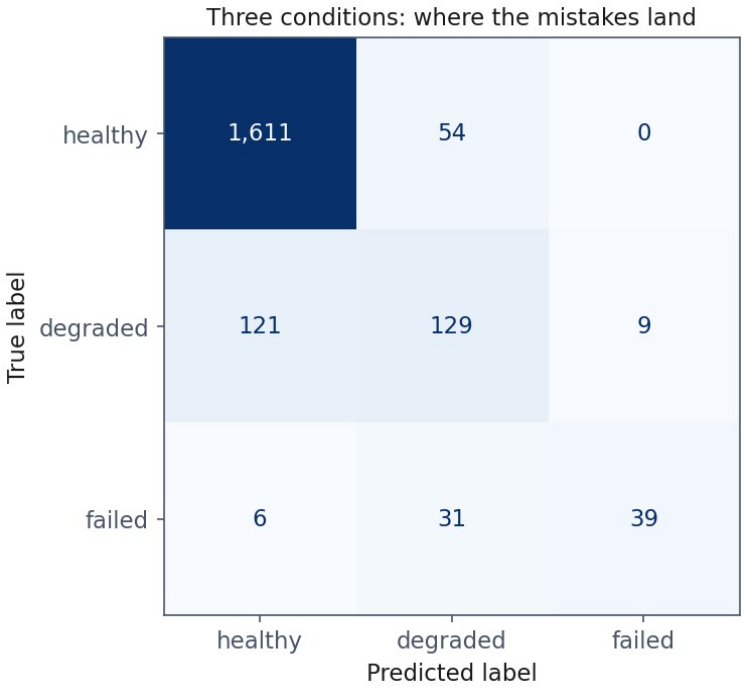
Model	Recall	Precision	AUC	Cost
Plain, threshold 0.50	0.566	0.768	0.949	87,620
Plain, threshold 0.08	0.855	0.349	0.949	45,540
Balanced class weights, 0.50	0.895	0.249	**0.950**	49,500

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Notebook 3, live

section, and its slides	you call (→ the number to come back with)
the ROC curve by hand, 37	`roc_curve`, `roc_auc_score` → 0.9493
what AUC means, 38	200,000 random (failing, healthy) pairs → 0.9488
the rare-positive trap, 40-41	`average_precision_score` beside `roc_auc_score` → AP 0.706 to 0.361, AUC unmoved
the threshold from costs, 43-45	`confusion_matrix` at each threshold → 45,540 EUR at 0.08, against 87,620 at 0.50
class weights, 46	`LogisticRegression` with balanced class weights → AUC 0.950
three classes, 48-49	`classification_report` → macro F1 0.707, weighted 0.883

More than two classes



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Macro against weighted, one last time

- macro avg F1: **0.707**
- weighted avg F1: **0.883**

UNIVERSITÀ
DEGLI STUDI
DELL'AQUILA

What to take away

- **96.20% accuracy, nothing caught.** Accuracy measures the majority class
- Precision is about the **alarms**, recall is about the **failures**
- The **threshold is a decision**, and it has a formula
- **AUC measures ranking** and hides imbalance

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Homework

- **Exercise 4**, discussed at the start of **Friday 30 October**
- Build a classifier, choose a threshold from a stated cost ratio, defend it

UNIVERSITÀ DEGLI STUDI DELL'AQUILA