

UNIVERSITÀ
DEGLI STUDI
DELL'AQUILA

Lesson 5: Experimental Methodology

Technologies for Artificial Intelligence

Fabio Antonini, Università degli Studi dell'Aquila

6 November 2026

We broke a promise on purpose

- Lesson 4 chose a threshold **by looking at the test set**
- Lesson 1 forbade exactly that
- Today we pay the debt

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Today: how to tell whether it works

- How far apart can two honest splits be?
- Cross-validation: what it estimates, and what it does not
- Bias, variance, and the error no model can remove
- Learning curves: more data, or a different model?
- Leakage that survives cross-validation
- Lesson 4’s threshold, done properly

UNIVERSITÀ DEGLI STUDI DELL’AQUILA

A test score is a measurement

- You would not report a poll of forty people to three decimals
- Holding data out makes the estimate **unbiased**: right on average over all the test sets you might have drawn, for the model you fitted
- It does nothing to make it **precise**: close to the truth on the one split you did draw

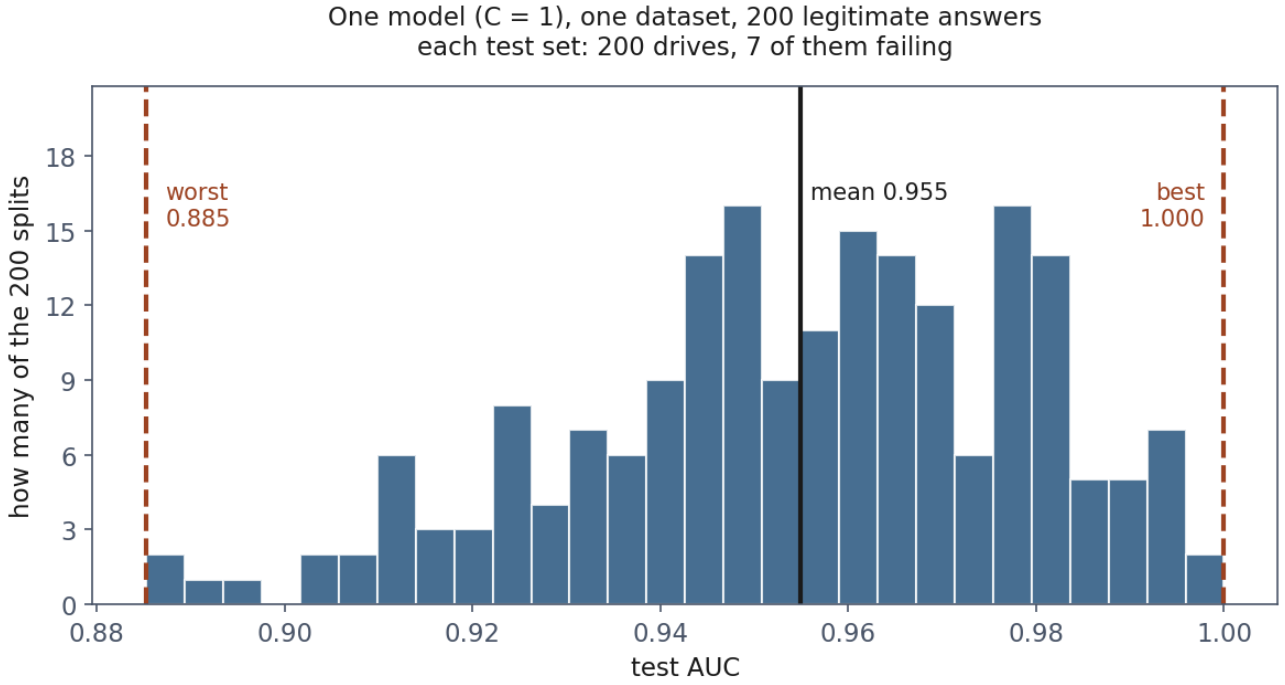
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The fleet, cut down

- Lesson 4’s drives, cut to **800**, of which **29 fail** (3.6%)
- Why 800: real projects often have a few hundred rows, a few dozen positive
- In code: `train_test_split(X, y, test_size=0.25, stratify=y)`; lesson 4’s pipeline fitted on the training part; the area under the curve (AUC), `roc_auc_score`, on the test part

UNIVERSITÀ DEGLI STUDI DELL’AQUILA

200 honest splits, one dataset



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Why do the 200 scores spread so widely?

- Each test set takes 25% of the 29 failures: 0.25 × 29 = 7.25, so **7 failures** among 200 drives
- AUC = pairs ranked right ÷ 1,351, the 7 × 193 = 1,351 pairs of one failing and one healthy drive
- For example, one failing drive slipping below 20 more healthy ones costs 20 ÷ 1,351 = 0.015 of AUC

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Parameters and hyperparameters

- **Parameters**: estimated by the fit, w and b
- **Hyperparameters**: not estimated by the fit; some fixed before it (penalty λ, neighbours, tree depth), some after it (the threshold)
- Training error would always pick $\lambda = 0$, so they are chosen by **comparing** models on data the fit did not see
- scikit-learn's C plays the role of $1/\lambda$: smaller C, stronger penalty

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Choosing a hyperparameter on a validation set: two recipes

Fixed before the fit: the penalty (C), tree depth, neighbours

1. Set 200 test drives aside; cut 150 validation drives from the other 600: 450 left to train

↓

2. For each candidate, C = 0.01, 1, 100: fit on the 450, score on the 150 = one fit each, 3 fits

↓

3. Keep the C with the best validation score: the winner is a value of C, not a model

↓

4. Refit with that C on all 600 non-test drives (fit number 4)

↓

5. Score the 200 test drives, once

Chosen after the fit: the decision threshold

1. Set 200 test drives aside; cut 150 validation drives from the other 600: 450 left to train

↓

2. Fit the model on the 450: one fit in all

↓

3. Predict probabilities on the 150 validation drives

↓

4. Try every threshold on those probabilities; keep the cheapest

↓

5. Score the 200 test drives, once

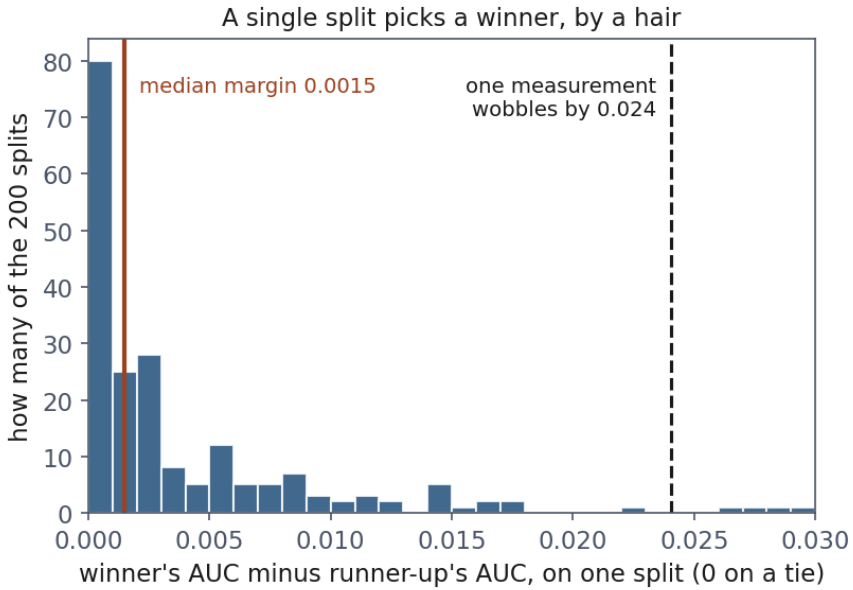
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

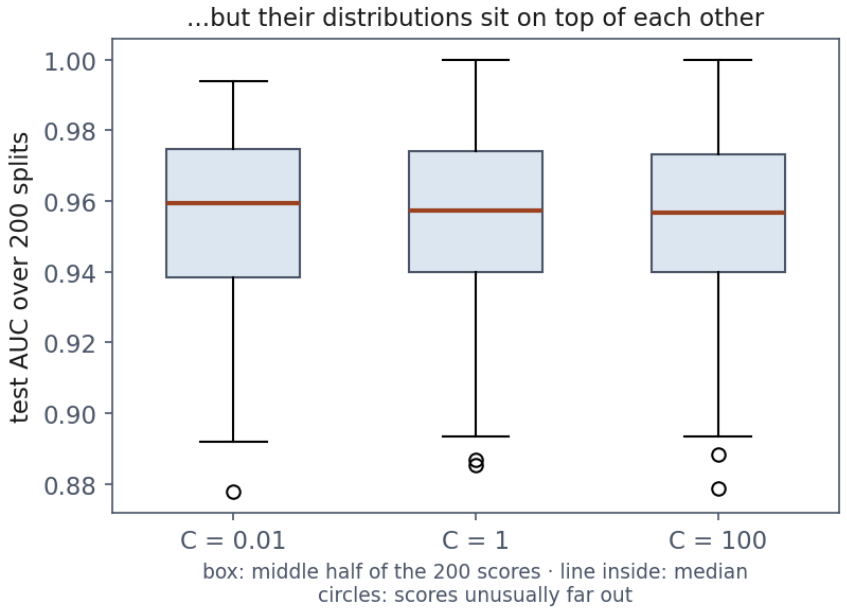
On each of slide 6’s 200 splits, three models compete

Model (C = 1, the default, is slide 6’s)	Mean test AUC over the 200 splits	Splits it won outright (31 were ties)
C = 0.01	0.9549	89
C = 1	0.9549	42
C = 100	0.9544	38

UNIVERSITÀ DEGLI STUDI DELL’AQUILA

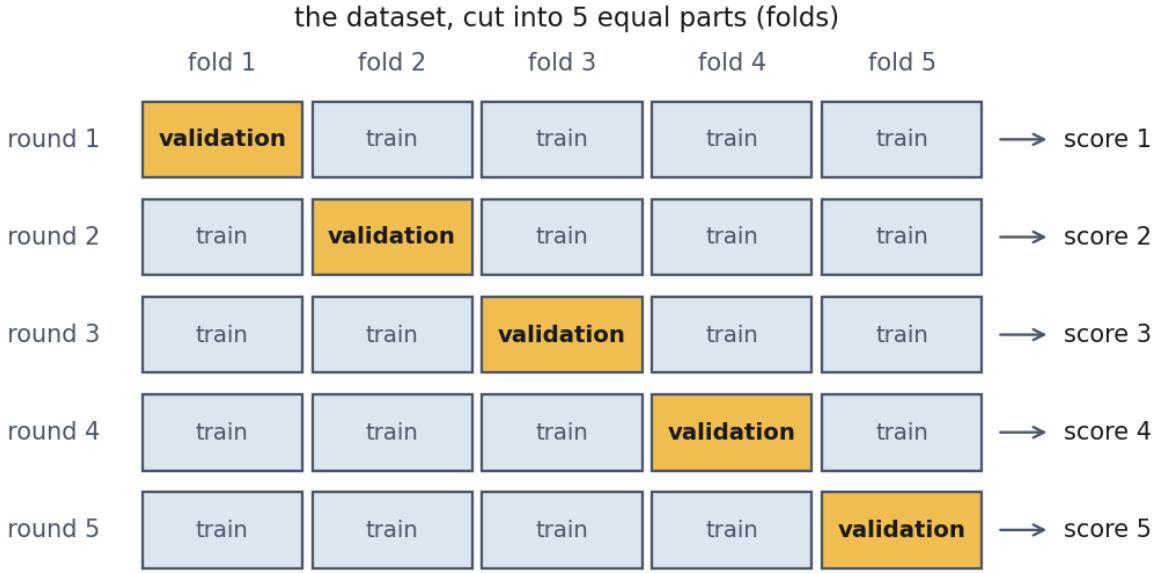
The same three models: who wins, and how they score





UNIVERSITÀ DEGLI STUDI DELL'AQUILA

k-fold cross-validation: the validation set, rotated



estimate =
(score 1 + ... + score 5) ÷ 5

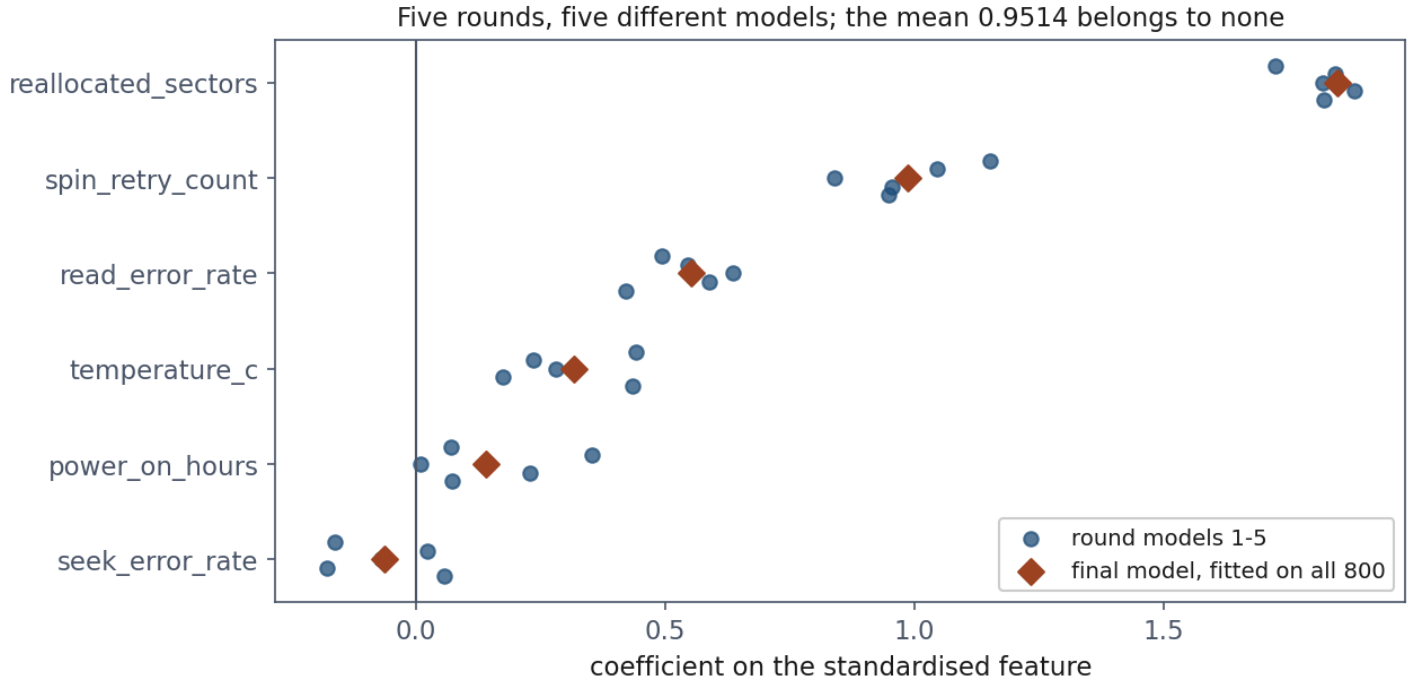
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

What it actually estimates

- **0.9514** (seed 0, mean of five round AUCs) scores a **procedure**, not one model
- Slightly **pessimistic**: each model trains on 80% of the data
- **Nothing chosen**: this is the measurement. **Something chosen**: test set aside first (slide 9)
- `cross_val_score` gives scores, no model; the model you ship is `fit` on all 800

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Five rounds, five different models; you ship a sixth, fitted on all 800

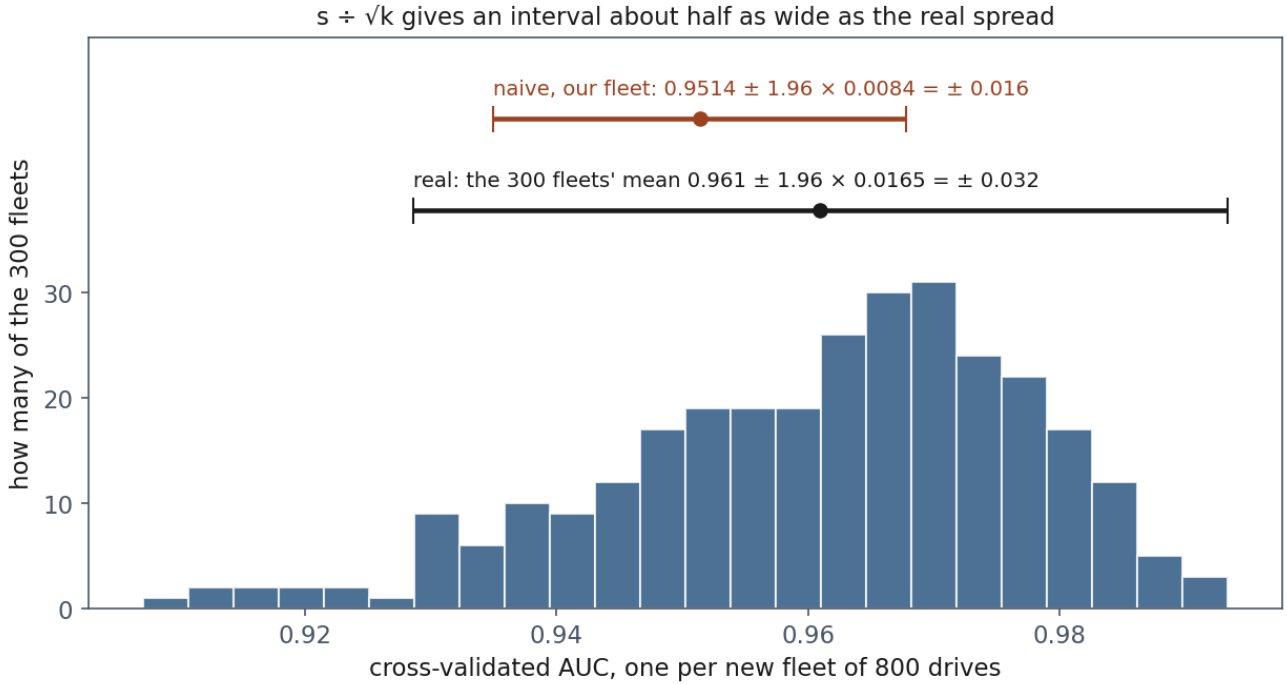


UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Report the spread, but not as a confidence interval

- On **another fleet of 800**, how far would 0.9514 move? A 95% **confidence interval** answers
- Tempting: ± 1.96 × s ÷ √5 = **± 0.016** (s = 0.0187: standard deviation of the 5 scores)
- **Too narrow**: the 5 scores share training data, and s rests on 5 numbers
- So: report mean and s, call s a **spread**, treat small differences as unresolved

300 new fleets: twice the naive spread; compare the widths



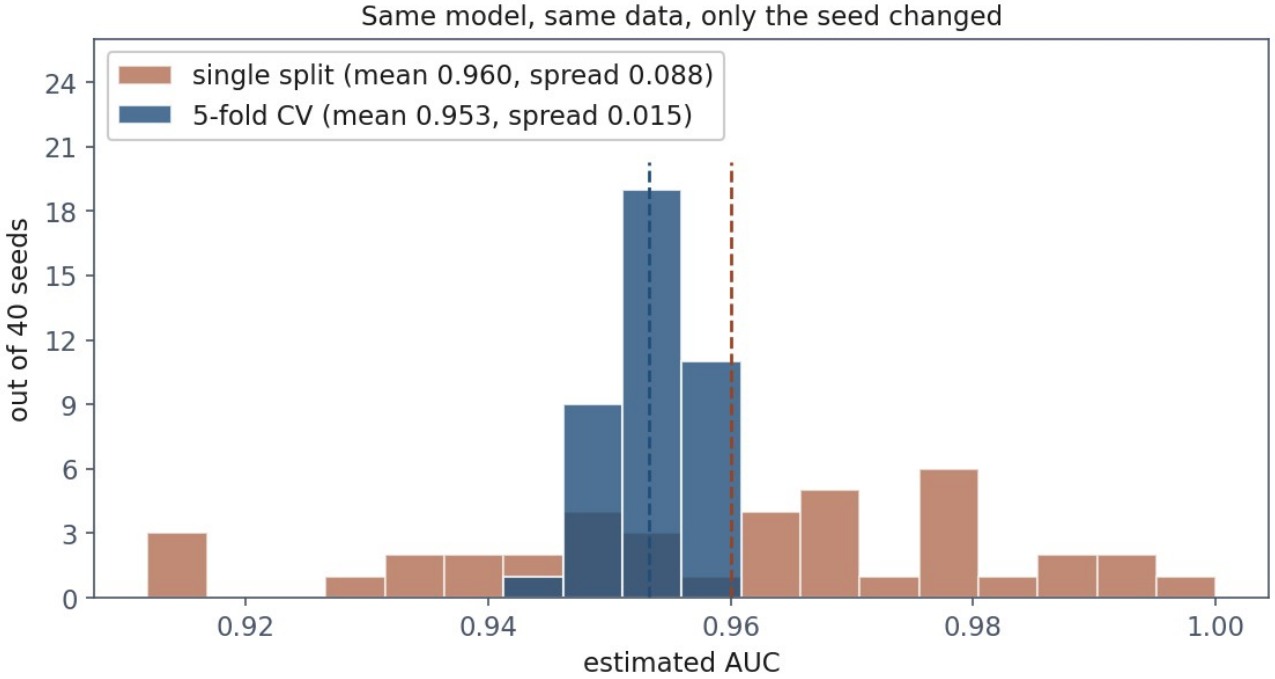
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Stratify when a class is rare

Fold (seed 3)	KFold: folds cut at random	StratifiedKFold: each fold keeps 3.6% failures
1	7 failures of 160, AUC 0.979	5 of 160, AUC 0.983
2	**1 failure of 160, AUC 1.000**	6 of 160, AUC 0.969
3	7 of 160, AUC 0.942	6 of 160, AUC 0.902
4	5 of 160, AUC 0.965	6 of 160, AUC 0.945
5	9 of 160, AUC 0.943	6 of 160, AUC 0.969

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Change only the seed, 40 times: how far does the score move? (optional)



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

How many folds? (optional)

- **k = 5 or 10** in practice
- Larger k: less pessimistic, more fits to run, round scores can be more correlated
- k = m (m the number of rows) is **leave-one-out**: for very few rows and a per-row metric; with AUC it returns nan (one row holds no pair)

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

When the rows are not independent

- Suppose ten readings per drive (lesson 4’s table has one)
- A random 5-fold split puts about 8 of a drive’s 10 readings in training, 2 in validation
- Those 2 are near-copies of the 8, same label: the model recognises the drive (1-nearest neighbour: 1.000 split at random, 0.620 by drive)

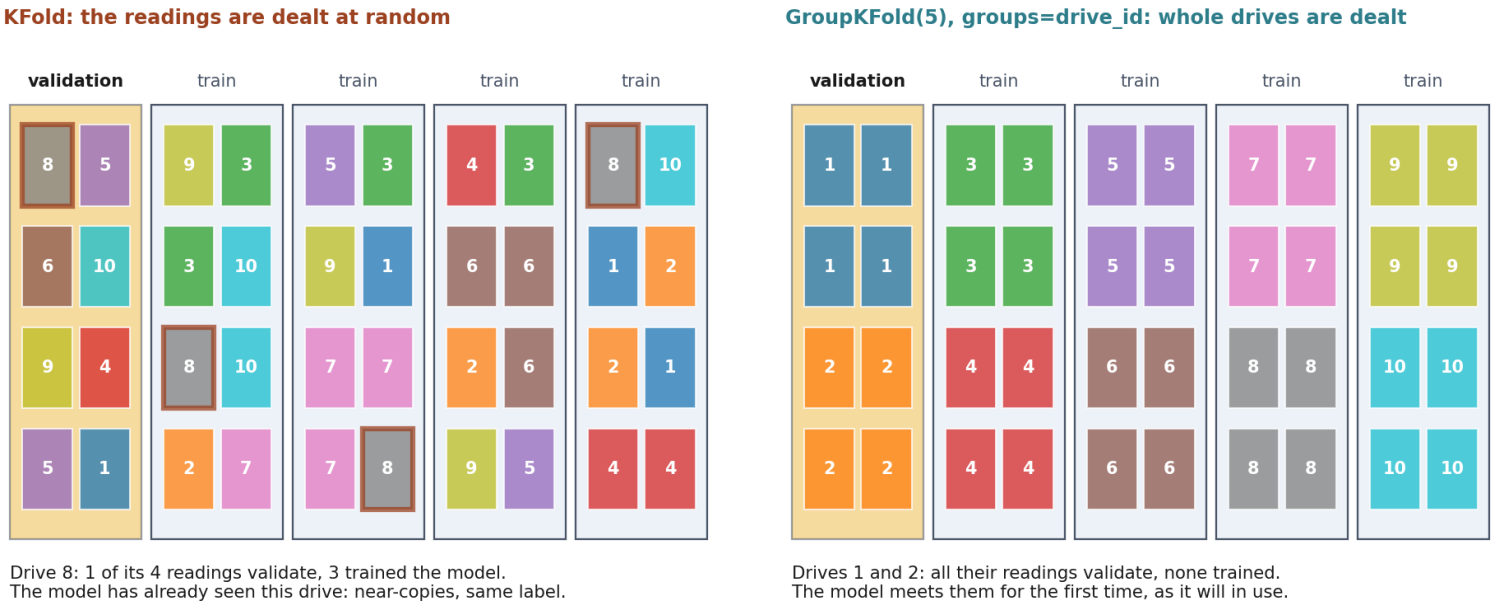
UNIVERSITÀ DEGLI STUDI DELL’AQUILA

Split by the thing you must generalise to

- `drive_id`: one value per row, the drive's serial number; not a feature, passed apart
- **GroupKFold** keeps a drive's rows in one fold: 640 drives train, 160 others validate
- `cross_val_score(model, X, y, groups=drive_id, cv=GroupKFold(5))`
- Test split too: `GroupShuffleSplit`; rare class: `StratifiedGroupKFold`

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

GroupKFold, drawn: each drive in one fold



Drawn small: 10 drives, 4 readings each, one square per reading, numbered by drive. On the lesson's data, 800 drives and 10 readings each:
1-nearest neighbour scores 1.000 split at random and 0.620 split by drive.

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Notebook 1, live

section, and its slides	you call (→ the number to come back with)
two honest experiments, 4-5	`train_test_split(..., random_state=seed)`, `make_pipeline(StandardScaler(), LogisticRegression())`, `roc_auc_score` → seed 3 gives 1.000
how wide, 6-7	the same, 200 times → 0.885 to 1.000
choosing on one split, 8-11	three values of C, a winner per split → 89, 42, 38 outright wins, 31 ties
k-fold by hand, 12-16	`StratifiedKFold`, then `cross_val_score` → mean 0.9514; 300 new fleets → 0.0165 against 0.0084
stratify, 17-18	`KFold` against `StratifiedKFold` → 1 failure, AUC 1.000; 0.0881 against 0.0149

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Break

- Ten minutes

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Error too high: more data, or another model?

- Daily energy use from the outdoor temperature, as in lesson 3
- A straight line on 25 days: 5-fold cross-validated mean squared error (MSE) **6,064.3 kWh²**, a typical miss of √6,064.3 ≈ 78 kWh
- **More days** or **another model**: each cures a different error
- Find where the error comes from: first in a lab, then on real data

A lab where we know the truth

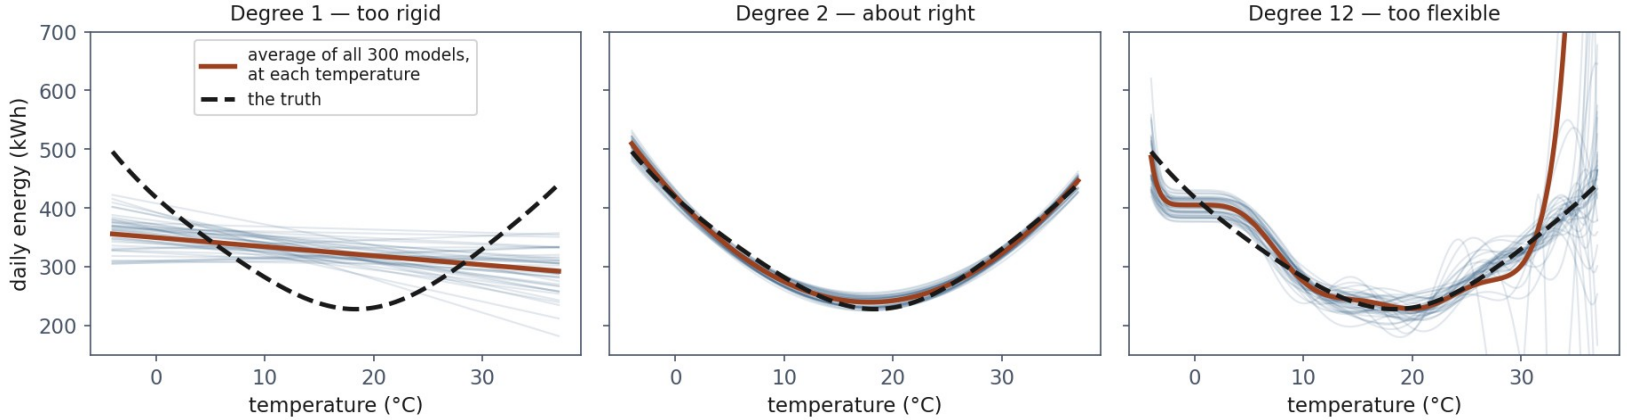
- Truth: $f(t) = 240 + 0.55\,(t - 18)^2 + 12\sin(t/4)$ kWh, t = daily mean temperature (°C); noise 22 kWh (standard deviation)
- 300 new training sets of 25 days at −5 to 38 °C, each drawn afresh (slide 25’s is the first)
- Polynomials of degree 1 to 12:

```
make_pipeline(PolynomialFeatures(d),
LinearRegression())
```


UNIVERSITÀ DEGLI STUDI DELL’AQUILA

Three hundred parallel universes: one model each

Each faint line is one model, fitted to its own 25 days (40 of the 300 drawn). Bias: the red average against the black truth. Variance: how widely the faint lines spread around the red average.



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The bias–variance decomposition: expected error = bias² + variance + noise

$$\mathrm{E}\left[(y - \hat{f}(x))^2\right] = \mathrm{bias}^2 + \mathrm{variance} + \sigma^2$$

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

What each term means

- A new day at x = 30 °C: y = its energy, truth plus fresh noise (kWh); $\hat{f}(x)$ = ONE model's prediction; $\mathbb{E}$ = average over models and noise
- **Bias** = average prediction − truth (kWh, signed); the mean squared error, expanded, yields bias²
- **Variance** = average of (prediction − average prediction)²: the models' disagreement
- **Noise** $\sigma^2 = 22^2 = 484$: the days' scatter; known because we added it

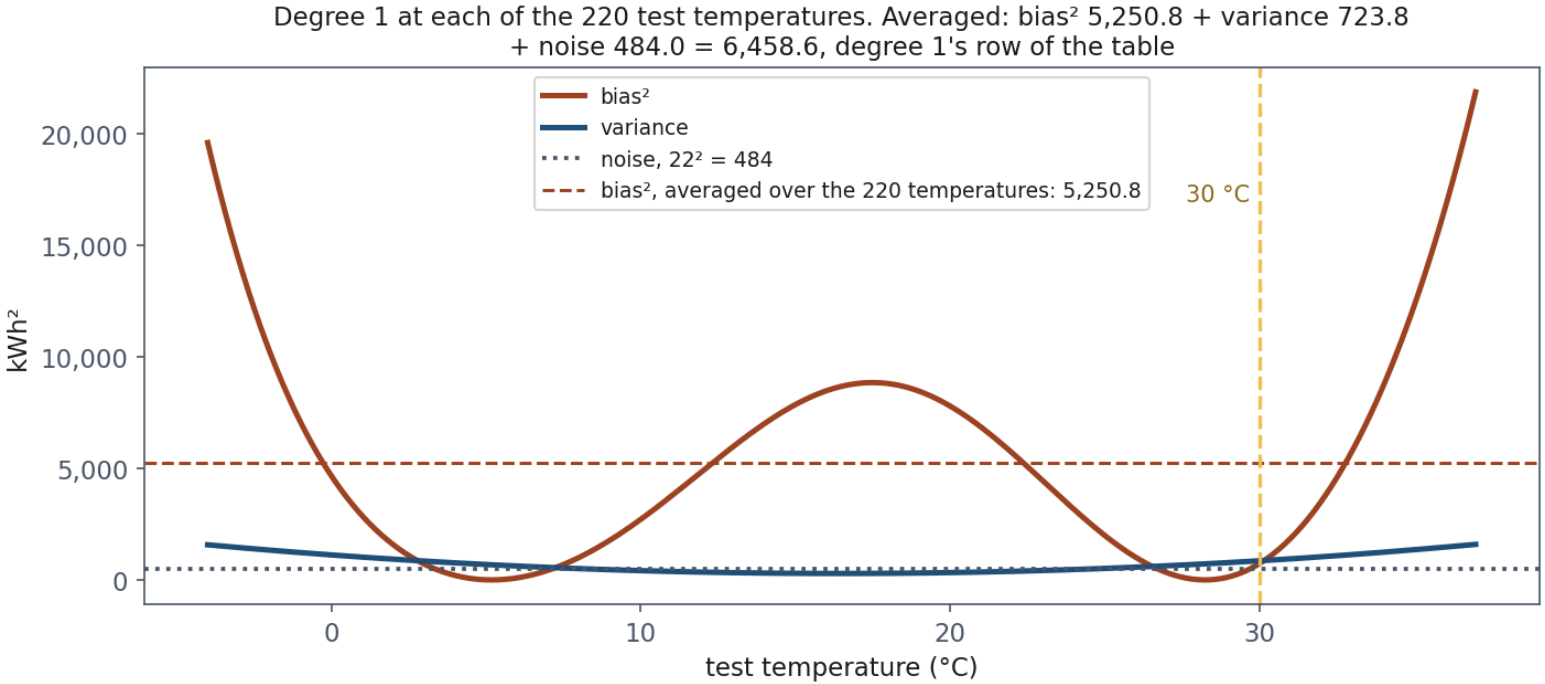
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The straight line (degree 1) at 30 °C, in three numbers

Step, degree 1	Formula	Value
The truth	f(30), known because we wrote f	330.5 kWh
The 300 predictions	one per model	300.8, 317.1, 304.8, and so on, from 215.6 to 364.7 kWh
Their average	sum ÷ 300	303.6 kWh
bias²	$(303.6 - 330.5)^2$	**724 kWh²**
Variance	sum of $(\text{prediction} - 303.6)^2$ ÷ 300	**856 kWh²**
Noise	22^2	**484 kWh²**
Expected error	bias^2 + variance + noise	**2,064 kWh²**

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Degree 1 at 220 temperatures, −4 to 37 °C, then averaged (optional)

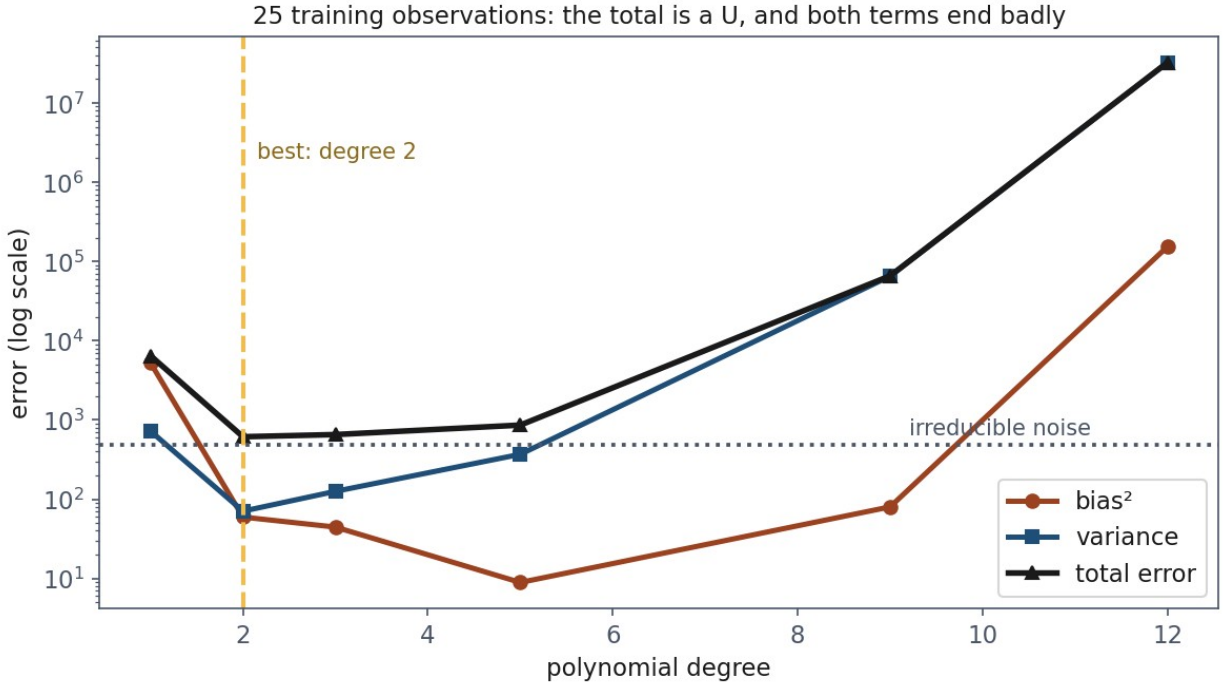


UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Every degree, on 25 days (optional)

Degree	bias² (kWh²)	Variance (kWh²)	Noise (kWh²)	Total = bias² + variance + noise (kWh²)
1	5,250.8	723.8	484.0	6,458.6
2	59.6	70.8	484.0	**614.3**
3	44.5	125.8	484.0	654.4
5	8.9	368.9	484.0	861.8
9	80.1	65,186.0	484.0	65,750.1
12	153,537.3	32,174,344.9	484.0	32,328,366.2

The bias–variance trade-off: the total error is a U



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The floor you cannot go below (optional)

- Noise variance is **484** = 22^2, the noise put into the data, at every complexity
- A **test** error well below a known floor means leakage or a mistake in the evaluation, not excellence

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The best model depends on how much data you have (optional)

Degree; total error at training-set size m (bold: lowest in its column)	m = 25	m = 60	m = 150
1	6,458.6	5,998.6	5,861.5
2	**614.3**	570.6	555.3
5	861.8	**540.9**	**506.9**
12	32,328,366.2	7,626.9	899.0

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Without the truth: two errors you can still measure

- One training set, f and σ unknown: cross-validation inside it, test set untouched, gives two errors
- **Training error**, on the days the model was fitted to: it has followed them, so it hides its variance
- **Validation error**, on the fold it did not see: it pays bias, variance and noise
- Training error too high for the task: **bias**. Validation far above training: **variance**

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

One sample of 25 days, 5-fold cross-validation

Degree	Training MSE, kWh²	Validation MSE, kWh²	Reading	What the lab measured
1	5,155.7	6,064.3	both high and close: **bias**	bias² 5,250.8
2	312.8	**450.4**	lowest validation: **the best**	lowest total, 614.3
5	262.4	615.1	training below the noise, 484; gap 615.1 − 262.4 ≈ 353: **variance**	variance 368.9 against degree 2's 70.8

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

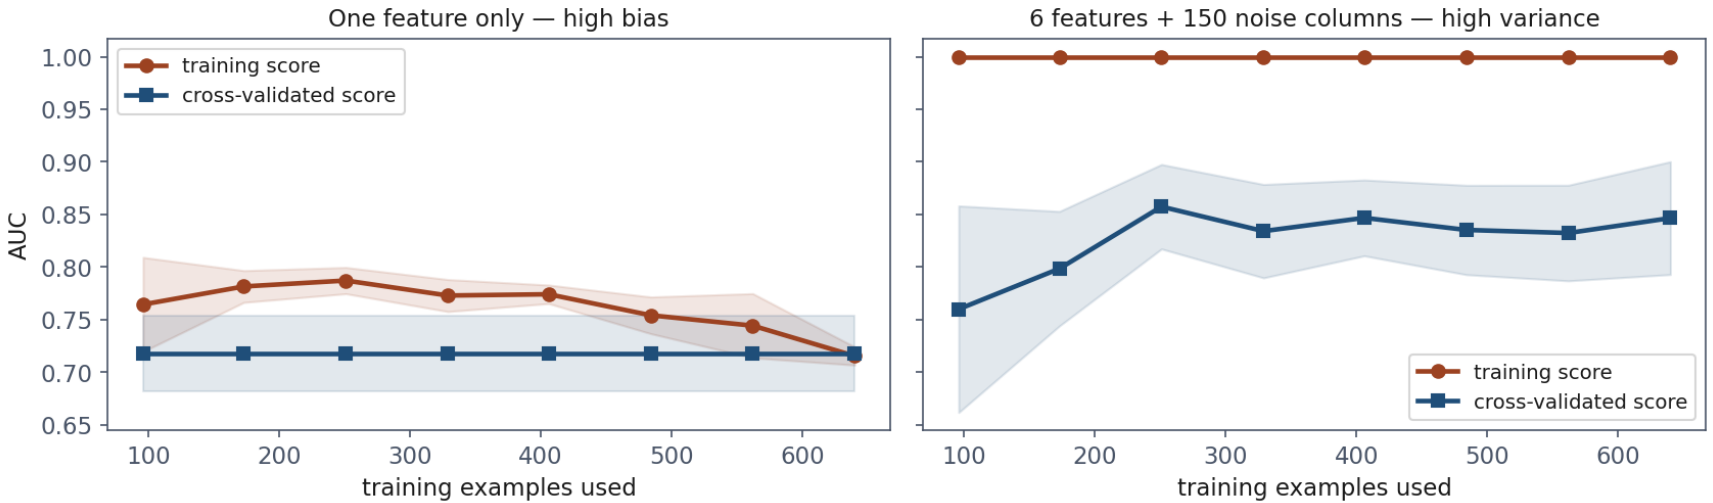
Learning curves: would more data help?

- Slide 37’s two errors, repeated at growing training-set sizes: `learning_curve(model, X, y, cv=folds, scoring="roc_auc")`
- On slide 5’s 800 drives (lesson 4’s fleet, cut down), scored by the area under the curve (AUC): **higher is better**, so poor = **low** curves
- Two models broken on purpose: one feature only; 6 features + 150 noise columns, almost no penalty

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Learning curves: left the curves meet low, right a gap that stops closing

Slide 5's 800 drives (lesson 4's fleet, cut down), logistic regression, 5-fold cross-validation.
The score is the AUC: higher is better, so a poor model shows as LOW curves.
Bias: the two curves meet low. Variance: a gap between them.



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Which fix to buy

- Curves meet at a **poor** score: bias; more data will **not** help
- Wide gap, validation **still rising** by more than its spread: more data **likely** helps
- Wide gap, validation **flat**: smaller C or **remove features**
- Slide 25’s answer: bias; degree 2, same 25 days: MSE **450.4**, not 6,064.3
- The test set: scored once, after the fix

Notebook 2, self-study (optional)

section, and its slides	you call (→ the number to come back with)
300 parallel universes, 26-27	`PolynomialFeatures` then `LinearRegression`, in `make_pipeline`, fitted 300 times
the three parts measured, 28-34	NumPy on the 300 predictions → sum and measured error agree to 2.73e-12
more data moves the answer, 35	the same at 25, 60 and 150 days → best degree 2, then 5
without the truth, 25 and 36-37	`cross_validate` with `KFold`, `return_train_score=True` → degree 1: 5,155.7 against 6,064.3
learning curves, 38-40	`learning_curve` with `StratifiedKFold` → 0.715 against 0.718; 1.000 against 0.847

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Now: what cross-validation does not protect you from

- Slide 5’s 800 drives and labels; features: **2,000 columns of random noise**; honest AUC **0.500**
- Keep the 10 columns with the highest F-test (failing against healthy means, relative to the spread in each group): `SelectKBest(f_classif, k=10)`
- Cross-validate a logistic regression on them: 5 rounds, train on 640, score 160

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The code almost everyone writes

```
model = make_pipeline(StandardScaler(),
LogisticRegression(max_iter=2000))
cv = StratifiedKFold(5, shuffle=True,
random_state=0)
sel = SelectKBest(f_classif, k=10)
sel.fit(noise, y)
X = noise.loc[:, sel.get_support()]
cross_val_score(model, X, y, cv=cv,
scoring="roc_auc")
```


UNIVERSITÀ DEGLI STUDI DELL'AQUILA

0.931

- Cross-validated AUC **0.931**, the mean of five folds: 0.928, 0.979, 0.975, 0.943, 0.832, on data with no signal
- The bug is `sel.fit(noise, y)`: the selector saw **every row**, validation folds included
- It kept the 10 columns that best matched the labels it would later be graded on

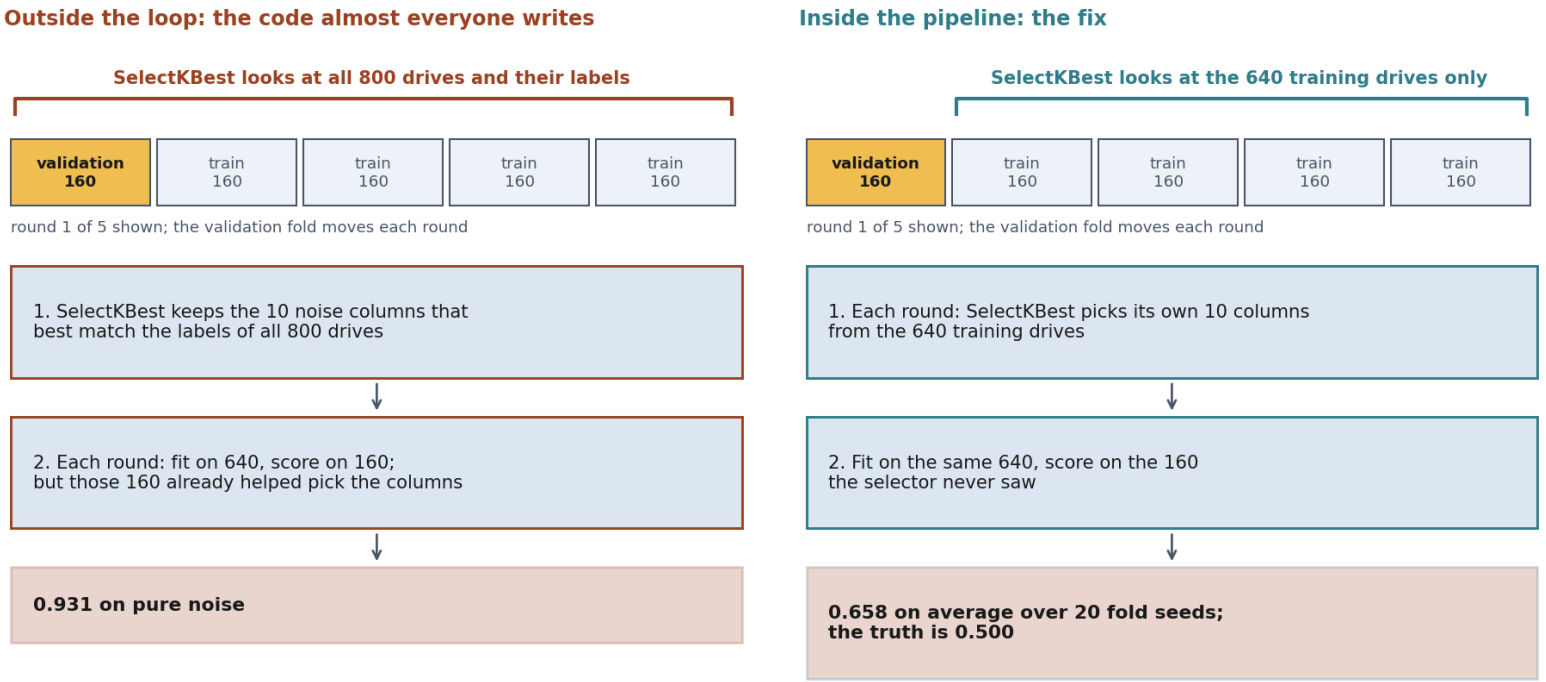
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Fix it, and look what happens

- Selection moved **inside** the pipeline, refitted in every round: `Pipeline([("select", SelectKBest(f_classif, k=10)), ("scale", StandardScaler()), ("model", LogisticRegression())])`
- 20 different fold seeds, 0 to 19: **0.658** on average, 0.576 to 0.759
- The truth is still 0.500

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The fix, drawn: the selector sees only training rows



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Why it does not return to 0.5

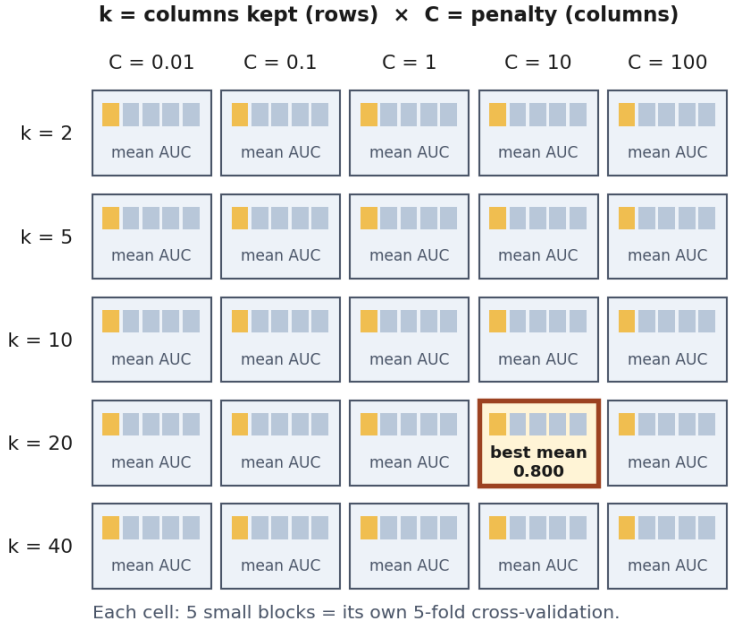
- Some of 2,000 columns match the label **by accident**: in the **sample**, so in every fold
- Helps: a test set set aside **before** any search; its chance matches sit in other columns
- Fewer candidates: the luckiest of 20 is less lucky than the luckiest of 2,000
- More rows: a noise column’s mean over 29 failures wobbles by 0.19; over 290, by 0.06

GridSearchCV: try all, keep the best

- **k** = columns kept: 2, 5, 10, 20, 40 (**columns, not folds**)
- **C** = penalty, smaller is stronger: 0.01, 0.1, 1, 10, 100
- 5 values of k × 5 values of C = **25 pairs**; each pair: 5-fold cross-validation, 25 × 5 = **125 fits**
- Best mean AUC wins: `best_params_` (k = 20, C = 10); `best_score_` = that mean
- Then refitted on all 800, **not scored again**

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

GridSearchCV, drawn: 25 cells, 5 folds each



1. For each of the 25 cells: 5-fold cross-validation, 5 fits each, 25 × 5 = 125 fits

↓

2. The cell with the highest mean AUC wins: best_params_ = k 20, C 10; best_score_ = 0.800

↓

3. Refit that cell's settings on all 800 drives (fit 126): the model you ship, not scored again

↓

0.800 is the best of 25 means, measured on the folds that also chose it; the truth is 0.500

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The best of 25 scores is inflated: 0.800 for a worthless model

`GridSearchCV`: k = 2, 5, 10, 20, 40 columns × C = 0.01, 0.1, 1, 10, 100	AUC
best of the 25 candidates, `best_score_`	**0.800**
average candidate	0.726
worst candidate	0.684
the luck of the winner: best − average = 0.800 − 0.726	0.074
the truth, by construction	0.500

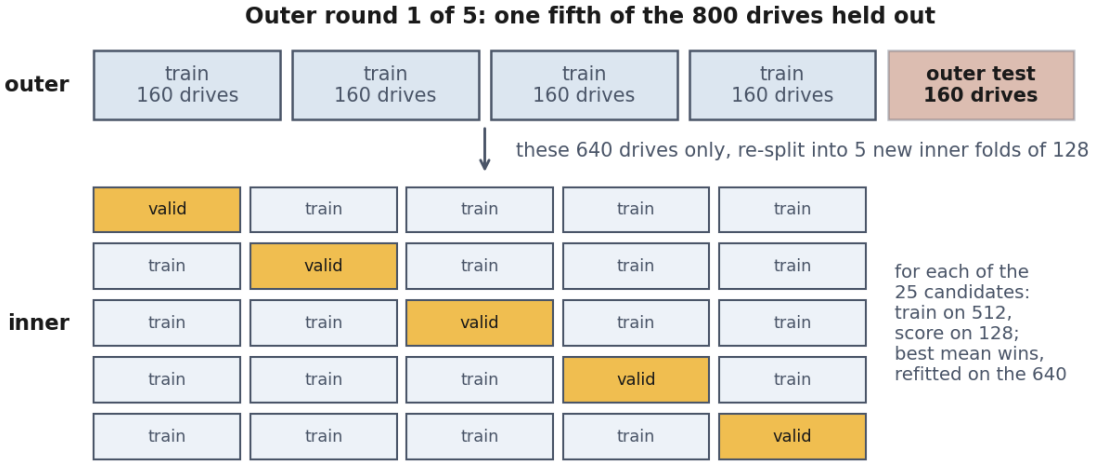
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Nested cross-validation: slide 9’s recipe, five times

- `GridSearchCV`’s folds both **choose and score**: `best_score_` is the luckiest of 25
- Set aside **160 drives** as the outer test
- Run the **whole search** on the other 640 (its own 5 inner folds), refit the winner on the 640
- Score it **once** on the 160; rotate, 5 times; mean of the 5 outer scores: **0.670**

UNIVERSITÀ DEGLI STUDI DELL’AQUILA

Nested cross-validation: it measures what choosing costs



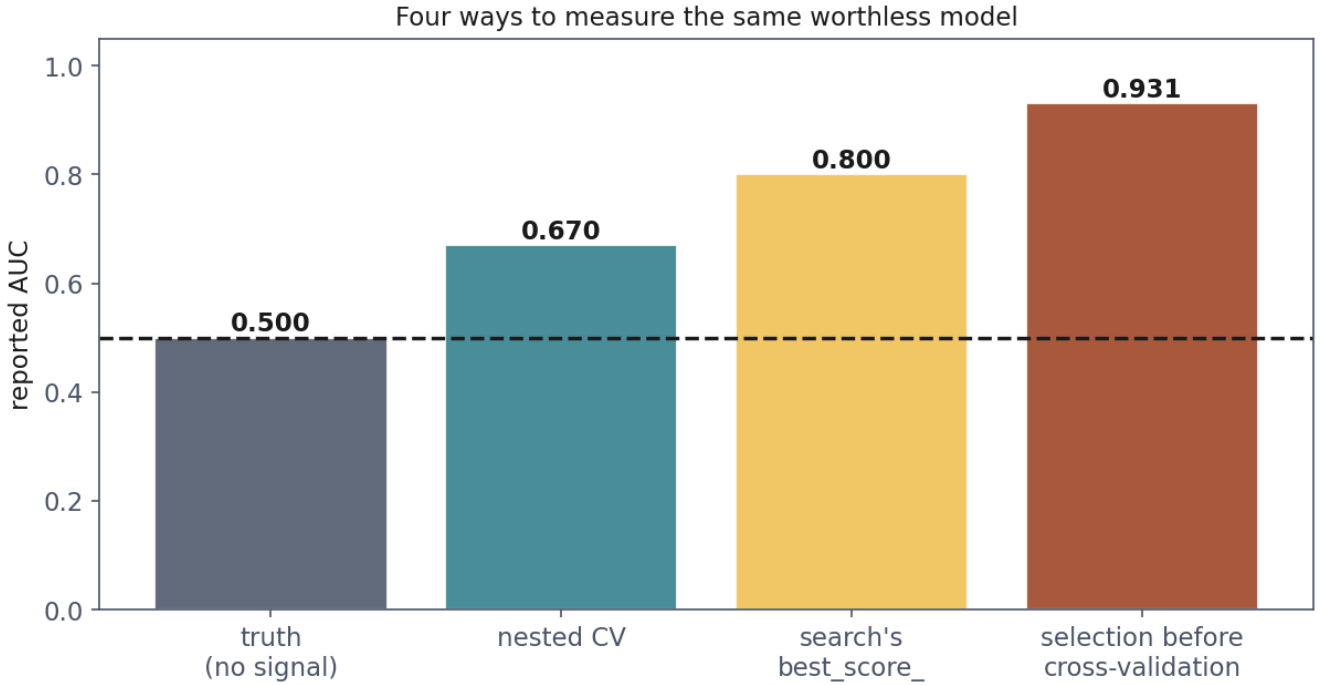
The winner is scored once on the 160 outer test drives: 0.698.

5 outer rounds: (0.698 + 0.852 + 0.754 + 0.487 + 0.558) ÷ 5 = 0.670

It measures what choosing costs; it does not choose.

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Each piece of discipline costs score: 0.931, 0.800, 0.670, 0.500



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The debt, paid: the honest threshold costs 1,400 EUR more

Threshold (one model, fitted on 450 drives, for all three rows)	missed failures × 2,600 EUR + false alarms × 140 EUR	Cost on the 200 test drives (7 failures)
0.50, the default: nothing chosen	4 × 2,600 + 1 × 140	10,540 EUR
0.050: the cheapest of 200 thresholds from 0.01 to 0.9, tried on the 150 **validation** drives	1 × 2,600 + 15 × 140	**4,700 EUR**
0.171: the same sweep, on the 200 **test** drives (lesson 4's way)	1 × 2,600 + 5 × 140	3,300 EUR

Three sets, three jobs

- **Train**: fit the parameters
- **Validation** (a held-out set, or the cross-validation folds): make every choice: model, hyperparameters, threshold
- **Test**: report. Touched **once**
- Change anything after looking at the test set, and it has become a validation set: **say so in the report**, or get fresh data

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

A fixed seed buys repeatability, not stability (optional)

- 30 seeds of the train/test split (600 drives to fit, 200 to test), same model: **0.913 to 1.000**
- Every run is perfectly reproducible
- `AUC 1.000 (random_state=3)` is reproducible **and** misleading
- Report the seed **and** the spread: AUC 0.951 ± 0.019 over 5-fold cross-validation, seed 0

What to fix, and what to write down (optional)

- Fix every `random_state` (the split, the folds), and seed your own NumPy draws with `np.random.default_rng(seed)`: a rerun gives the same number
- **Library versions**: defaults change between releases
- What you could **not** fix: parallel computations that finish in a different order, graphics processing unit (GPU) non-determinism

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Notebook 3, live

section, and its slides	you call (→ the number to come back with)
the catastrophe, 42-44	`SelectKBest(f_classif, k=10)` on every row, then `cross_val_score` → 0.931
selection inside the round, 45-47	`Pipeline` with the selector as its first step → 0.658
the search, 48-50	`GridSearchCV`, `best_score_` → 0.800
nested, 51-53	`cross_val_score(GridSearchCV(honest, grid, cv=inner), noise, y, cv=outer)` → 0.670
the threshold, honestly, 54-55	`train_test_split` twice → 4,700 EUR against 10,540
seeds, 56-57	30 values of `random_state` → 0.913 to 1.000

What to put in a report

- The metric, its **spread** (standard deviation of the fold scores, slide 15), and how it was estimated (e.g. stratified 5-fold, seed 0)
- How many **positives** the test set held
- Every seed, and every library version
- Which choices were made on which data

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

What to take away

- **A score is a measurement**: 0.885 to **1.000**
- **Choosing on one split chooses noise**
- **High training error: bias; a gap: variance**
- **Choose on validation, report on test once**
- **What cross-validation cannot remove**:
  - chance correlations in the sample: 0.658 on pure noise
  - the luck of choosing: 0.800 reported, 0.670 nested
  - dependent rows, unless split with `GroupKFold`

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Homework

- **Exercise 5**, discussed at the start of **Friday 20 November**
- A result that is too good. Find out why, then produce the honest number