



# Lesson 6: k-NN, Naive Bayes and Support Vector Machines

Technologies for Artificial Intelligence

Fabio Antonini, Università degli Studi  
dell'Aquila

# Before we start

- Exercise 5, which we discuss now
- The score that was too good, and what the honest number cost you


# Today: the first lesson that offers a choice

- **k-nearest neighbours (k-NN)**: remember, and vote
- The **curse of dimensionality**: the bill for that
- **Naive Bayes**: one strong assumption, bought cheap
- **Support vector machines (SVM)**: margins, and kernels
- Which family to reach for, and why that beats tuning

# 1,200 pumps, two readings each

- Vibration in Hz, pressure in bar
- A pump is faulty when its readings fall **outside the design envelope**, too low as readily as too high
- **61.3%** faulty; 4% of the labels are deliberately flipped

# The shape of the problem



# What a straight boundary costs

| Model                             | Cross-validated accuracy            |
|-----------------------------------|-------------------------------------|
| Always predict the majority class | 0.613                               |
| Logistic regression (lesson 4)    | <b>0.613 <math>\pm</math> 0.000</b> |

# k-nearest neighbours: the whole algorithm

- To classify a new point, find the **k closest** training points and take the majority vote
- There is no training step: “fitting” means **storing the data**
- Called a **lazy learner**: an unusually honest name for an algorithm

Distance is all it has

$$d(x, x') = \sqrt{\sum_{j=1}^n (x_j - x'_j)^2}$$

# Scale first, or vibration decides everything

- Vibration runs over **tens of Hz**; pressure over **a couple of bar**
- Unscaled, a one-bar pressure difference is invisible beside a ten-Hz one
- Lesson 2's scaling argument, arriving with **immediate consequences**


# k is the bias-variance dial, made visible

- **Small k**: the boundary follows every point, mislabelled ones included. Low bias, high variance
- **Large k**: the vote is taken over a wide neighbourhood, and eventually stops following real structure. High bias, low variance
- At  $k = 1$  the training accuracy is **exactly 1.000**. Always, on any dataset


# Measured on the pumps

| k   | Training     | Cross-validated |
|-----|--------------|-----------------|
| 1   | <b>1.000</b> | 0.912           |
| 5   | 0.953        | <b>0.944</b>    |
| 15  | 0.948        | 0.938           |
| 51  | 0.940        | 0.933           |
| 401 | 0.828        | 0.708           |

# The two curves, and the two lines that bound them



# The same data at three values of k



# What k-NN costs

- No training time at all: you pay at **every prediction** instead
- A naive implementation compares the query against **every** training row:  $O(mn)$  for  $m$  rows and  $n$  features
- **The model is the dataset.** You cannot ship one without the other

# Now the price: add columns of nothing

- Extra columns of **pure noise**, drawn from a normal distribution
- The two real readings are **untouched**: the problem is exactly as solvable
- Only the number of columns changes


# The signal never left

| Total columns | Accuracy | Above baseline |
|---------------|----------|----------------|
| 2             | 0.938    | +0.325         |
| 12            | 0.762    | +0.149         |
| 27            | 0.662    | +0.049         |
| 52            | 0.602    | <b>-0.011</b>  |
| 102           | 0.578    | -0.035         |

# Why: distances stop varying

- Scatter points in a cube, pick one, measure its **nearest** and its **farthest** neighbour
- In two dimensions those are very different numbers, which is what makes “nearest” a meaningful word
- Each new dimension adds its own squared difference, and the contributions **average out**

# The nearest point stops being near



# Nearest divided by farthest

| Dimensions | nearest ÷ farthest |
|------------|--------------------|
| 2          | 0.016              |
| 10         | 0.263              |
| 50         | 0.592              |
| 100        | <b>0.701</b>       |
| 500        | 0.855              |

# In 100 dimensions the nearest point is 70% as far as the farthest

- In **two** dimensions it is **2%**: “nearest” picks out something special
- In **one hundred** it is **70%**: the nearest point is barely nearer than a random one
- A vote among “the five nearest” becomes a vote among five taken at random

# What the curse is, and is not

- **Not** that high-dimensional problems are unlearnable: lesson 9's networks work in thousands of dimensions
- That **methods built on distance lose their footing**, because the quantity they depend on stops varying
- k-NN, k-means (lesson 8), and radial basis function (RBF) kernels, all of them

# Notebook 1, live

- k-NN on the pumps, choosing k, and the curse measured

# Break

- Twelve minutes

# Naive Bayes: turn the question around

- We want the probability of the class **given** the readings
- The prior is a count; the denominator is identical across classes and cancels
- Everything hard sits in one term: the probability of **this exact combination of readings** among examples of that class

# Bayes' rule, applied to a class

$$P(y = c \mid x) = \frac{P(x \mid y = c) P(y = c)}{P(x)}$$

The decision drops the denominator

$$\hat{y} = \arg \max_c P(x | y = c) P(y = c)$$

# The assumption: independent, given the class

- **Given the class, the features are independent of one another**
- The joint density then factorises into  $n$  one-dimensional densities
- Each factor is estimated from the rows of that class alone

The factorisation the assumption  
buys

$$P(x \mid y = c) = \prod_{j=1}^n P(x_j \mid y = c)$$



Gaussian: each factor is a bell curve

$$P(x_j \mid y = c) = \frac{1}{\sqrt{2\pi\sigma_{jc}^2}} \exp \left[ -\frac{(x_j - \mu_{jc})^2}{2\sigma_{jc}^2} \right]$$

# Ten numbers, and the widths decide

| Class   | Prior  | Vibration mean | Vibration sd | Pressure mean | Pressure sd  |
|---------|--------|----------------|--------------|---------------|--------------|
| Healthy | 0.3875 | 41.986 Hz      | <b>1.695</b> | 5.617 bar     | <b>0.219</b> |
| Faulty  | 0.6125 | 42.015 Hz      | <b>4.636</b> | 5.616 bar     | <b>0.576</b> |

# 48 Hz: one bell is 86 times the other



# 48 Hz, 5.6 bar: faulty, 0.981

| Class   | log prior | log P(vibration) | log P(pressure) | Total         |
|---------|-----------|------------------|-----------------|---------------|
| Healthy | -0.948    | -7.745           | +0.598          | -8.095        |
| Faulty  | -0.490    | -3.286           | -0.368          | <b>-4.144</b> |

# From the two totals to 0.981

- Each total is a log score:  $\log \text{prior} + \log P(\text{vibration}) + \log P(\text{pressure})$
- The gap:  $-4.144 - (-8.095) = \mathbf{3.951}$ , in favour of faulty
- $P(\text{faulty}) = 1 / (1 + e^{(-3.951)}) = 1 / 1.0192 = \mathbf{0.981}$
- Check:  $e^{(-4.144)} / (e^{(-4.144)} + e^{(-8.095)}) = 0.01585 / 0.01616 = 0.981$

# A wonderful bargain, and almost never true

- Training is **a single pass**; adding features costs almost nothing
- But vibration and pressure are both driven by the operating point
- “New” is not independent of “York” given the topic. Symptoms co-occur
- So the question is not whether it holds: it is **when being wrong about it costs you nothing**


# On the pumps, it scores 0.933

| Model                       | Accuracy     |
|-----------------------------|--------------|
| Majority baseline           | 0.613        |
| Logistic regression         | 0.613        |
| <b>Gaussian Naive Bayes</b> | <b>0.933</b> |
| k-NN, $k = 5$               | 0.944        |

# Uncorrelated, but not independent

| Within class  | Readings | Squared distances |
|---------------|----------|-------------------|
| healthy pumps | -0.006   | +0.168            |
| faulty pumps  | -0.049   | <b>-0.420</b>     |




# Wrong densities, right boundary



# One step away: when the signal is an interaction

- A second pair of sensors on the same fleet
- Faulty when **exactly one** of the two readings is high
- Both high is the designed high-load mode; both low is idle; one without the other is a mismatch between demand and delivery

# What Naive Bayes gets to see



# 0.404: below the baseline, and below chance

| Model                       | Accuracy     |
|-----------------------------|--------------|
| Majority baseline           | 0.523        |
| <b>Gaussian Naive Bayes</b> | <b>0.404</b> |
| Logistic regression         | 0.393        |
| SVM, linear kernel          | 0.606        |
| k-NN, k = 5                 | 0.967        |
| SVM, RBF kernel             | 0.972        |

# Its probabilities are not probabilities

- Mean confidence when **correct: 0.567**. When **wrong: 0.555**
- It cannot tell the difference between the two situations
- With features that repeat each other it fails the other way: 0.999 reported, with an accuracy nothing like that
- **The ranking may be useful while the probabilities are not**


# Notebook 2, live

- Where breaking the assumption is cheap, and one step away where it is fatal

# The margin: a criterion that is not the error

- On separable data there are **infinitely many** separating lines, and every one has **zero training error**
- Minimising the error therefore cannot choose between them
- The SVM picks the widest **slab**: push it out until it touches the nearest point of each class
- A boundary passing close to a point is one small perturbation from getting it wrong

# Three points decide everything



# Not separable? Charge for the violations

- Our labels are 4% flipped, so the strict problem has **no solution at all**
- Let each point violate the margin by  $\xi_i$ , and charge  $C$  for every unit of it
- $\xi_i = 0$  outside the slab, between 0 and 1 inside it, above 1 on the wrong side
- That is the **hinge loss**: exactly zero for a point with room to spare

# The hinge stops at zero



A wide slab, minus what the violations cost

$$\min_{w, b, \xi} \frac{1}{2} \|w\|^2 + C \sum_i \xi_i$$

# C is the price of a training error

- **Large C**: violations are expensive, so the model contorts to classify everything. Narrow margin, low bias, high variance
- **Small C**: a wider, calmer boundary, at the cost of some errors
- The same dial as  $k$  in  $k$ -NN and  $\lambda$  in lesson 3, in a third costume
- Note the direction: **large C means less regularisation**

# The best straight line is still a straight line

| Model              | Accuracy                            | Support vectors             |
|--------------------|-------------------------------------|-----------------------------|
| SVM, linear kernel | <b>0.613 <math>\pm</math> 0.000</b> | 947 of 1,200 ( <b>79%</b> ) |
| SVM, RBF kernel    | 0.947 $\pm$ 0.005                   | 278 of 1,200 (23%)          |


# A high support-vector fraction is a free warning

- The linear model needs **79%** of the training set to define its boundary
- The RBF model needs **23%**
- Almost every point sits on or inside the margin: there is no slab that separates anything
- The count falls out of `fit` at no cost, and you should look at it


# The kernel trick: they needed different coordinates

- Healthy pumps in a disc, faulty ones around it: no line works **in the plane**
- Add a third coordinate: the **distance from the design point**
- Healthy pumps rise a little, faulty ones a lot, and a flat plane separates them perfectly
- The classes were always separable. They were measured in the wrong coordinates

# The same pumps, lifted



Lifted: a flat plane does  
height = distance from the design point



# And the map is never computed

- The solution depends on the data **only through inner products** between pairs of points
- So we need a function returning the inner product **in the new space**, while computing only with the original one
- The radial basis function is scikit-learn's default: one exponential per pair, for an **infinite-dimensional** space

# The kernel trick by hand: $25 = 25$

| Route               | Computation                                    | Result    |
|---------------------|------------------------------------------------|-----------|
| Lift, then multiply | $(1, 2.83, 4) \cdot (9, 4.24, 1) = 9 + 12 + 4$ | <b>25</b> |
| Never lift          | $((1, 2) \cdot (3, 1))^2 = 5^2$                | <b>25</b> |

The radial basis function, written out

$$K(x, x') = \exp(-\gamma \|x - x'\|^2)$$


# Training: one weight per pump

- Weights adjusted **two at a time**; no learning rate, same answer every run
- **922** pumps: weight 0, forgotten
- **268** at the ceiling  $C = 1$ , **10** on the edge: the 278 support vectors
- The model: 278 points, their weights, and  $b = +1.105$

# The SVM on 48 Hz, 5.6 bar: +1.883

| Term                               | Value         |
|------------------------------------|---------------|
| votes from faulty support vectors  | +37.455       |
| votes from healthy support vectors | -36.677       |
| intercept b                        | +1.105        |
| <b>total</b>                       | <b>+1.883</b> |


# One pump's vote, drawn






# The highest training score, the lowest honest one

| Setting                 | Training     | Cross-validated |
|-------------------------|--------------|-----------------|
| $\gamma = 0.1, C = 1$   | 0.936        | 0.929           |
| $\gamma = 1, C = 1$     | 0.950        | <b>0.944</b>    |
| $\gamma = 50, C = 1000$ | <b>0.995</b> | 0.902           |

# How far one vote reaches



# Overfitting you can see



# Notebook 3, live

- Margins, support vectors, and what  $\gamma$  does to the boundary

# The three compared, on identical data

| Model                          | Accuracy       |
|--------------------------------|----------------|
| Majority baseline              | 0.613          |
| Logistic regression (lesson 4) | 0.613          |
| SVM, linear kernel             | 0.613          |
| Gaussian Naive Bayes           | 0.933          |
| k-NN, $k = 5$                  | 0.944          |
| SVM, RBF kernel                | <b>0.947</b>   |
| <i>noise ceiling</i>           | $\approx 0.96$ |

# How to choose, in practice

| Situation                                         | Reach for              |
|---------------------------------------------------|------------------------|
| Few features, plenty of data, odd-shaped boundary | k-NN, or an RBF SVM    |
| Very many features, little data                   | Naive Bayes            |
| Need a calibrated probability                     | <b>Not</b> Naive Bayes |
| Model must be small, or fast at prediction        | SVM, not k-NN          |
| The signal is an interaction between features     | k-NN or a kernel       |

# What to take away

- **No line separates the pumps:** 0.613, the base rate
- **k is the bias-variance dial**, made visible
- **In 100 dimensions the nearest point is 70% as far as the farthest**
- **Naive Bayes:** 0.933 where independence holds, **0.404** where it does not
- **Choosing the family beats tuning the wrong one:** 0.613 to 0.947

# Homework

- **Exercise 6**, discussed at the start of **Friday 13 November**
- A dataset, three families, and a defence of which one you chose