

UNIVERSITÀ
DEGLI STUDI
DELL'AQUILA

Lesson 6: k-NN, Naive Bayes and Support Vector Machines

Technologies for Artificial Intelligence

Fabio Antonini, Università degli Studi dell'Aquila

20 November 2026

Before we start

- Exercise 5: the colleague’s gearbox model looked excellent. What made its score too good, and what honest number did you report instead?
- How much did honesty cost, and which fix cost the most?

Today: the first lesson that offers a choice

- **k-nearest neighbours (k-NN)**, and the **curse of dimensionality** that limits it
- **Naive Bayes**: a probability model, very cheap to train
- **Support vector machines (SVM)**: the widest margin, then kernels
- How to choose a model family, and why that matters more than tuning

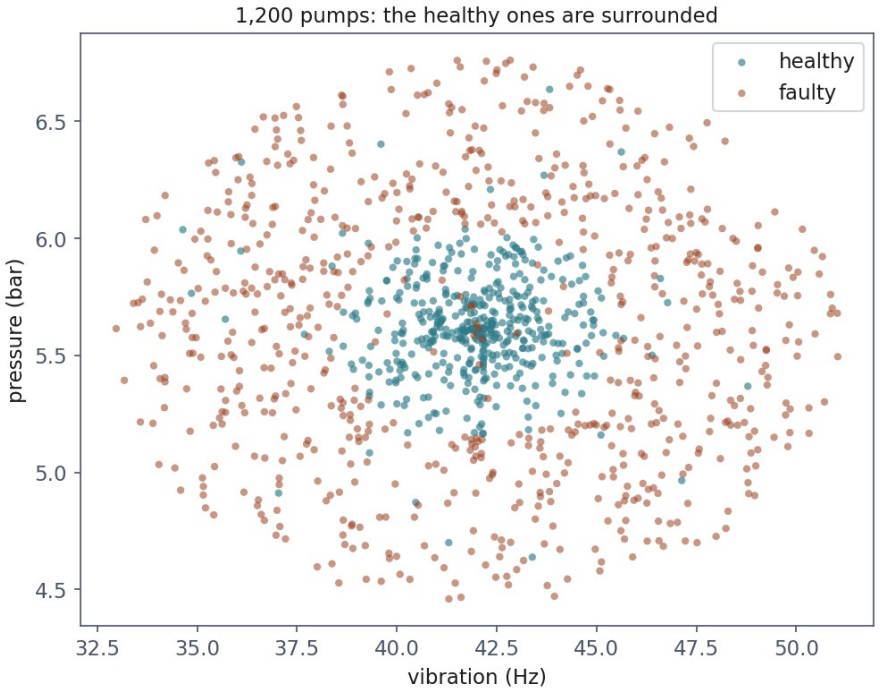
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

1,200 pumps, two readings each

- 1,200 simulated pumps: vibration in hertz (Hz), pressure in bar
- Faulty when, in steps of 3.5 Hz and 0.45 bar, the two readings together are more than one step from the **design point, 42 Hz and 5.6 bar**
- 735 of the 1,200 are faulty: **61.3%**
- 4% of the labels flipped on purpose: on new pumps, about **0.96** at best (the noise ceiling)

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The shape of the problem: healthy pumps in the middle, faulty all around



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

A straight boundary calls all 1,200 pumps faulty

Model	Accuracy, 5-fold cross-validation (mean ± standard deviation over the folds)
Always predict the majority class: 735 ÷ 1,200	0.613
Logistic regression (lesson 4)	**0.613 ± 0.000**

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

k-nearest neighbours: the whole algorithm

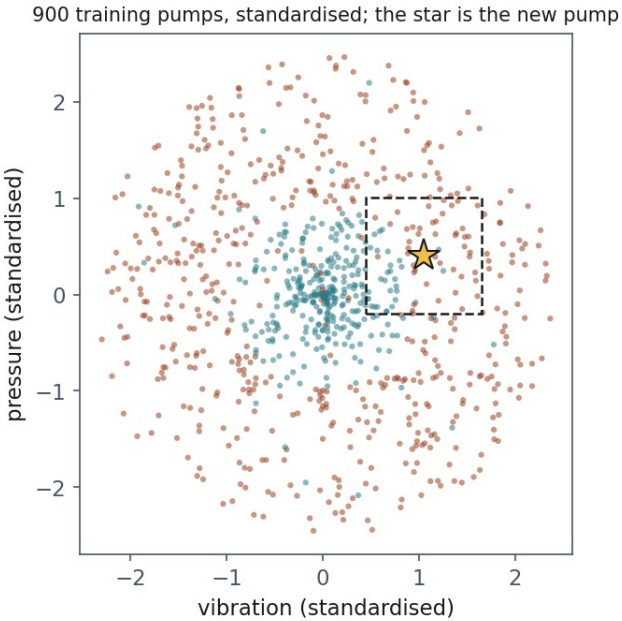
- A line fails, so try a method that assumes no shape: to classify a new pump, take the majority vote of the **k closest** training pumps
- There is no training step: “fitting” only **stores the data** (a **lazy learner**: the work waits until a prediction is asked for)
- In code:

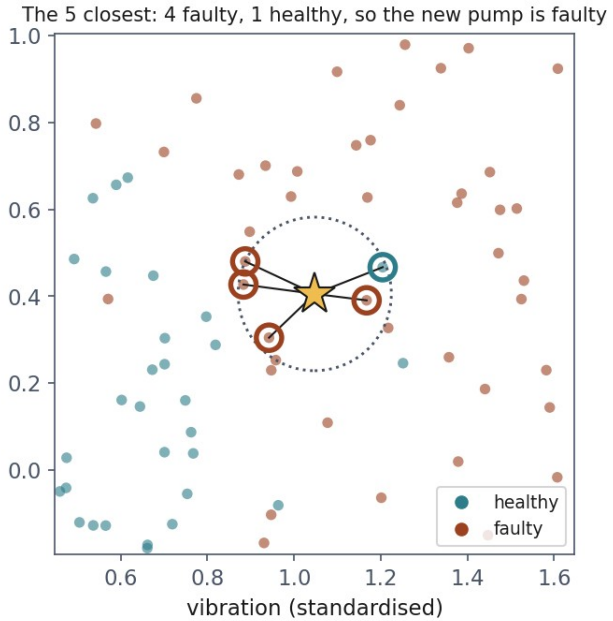
```
KNeighborsClassifier(n_neighbors=5)
```


UNIVERSITÀ DEGLI STUDI DELL'AQUILA

k-NN, drawn: one new pump and its five nearest

k-NN with k = 5: 1. standardise 2. measure the distance to every training pump 3. keep the 5 closest 4. majority vote





UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Euclidean distance between two pumps x and x′, over their n readings

$$d(x, x') = \sqrt{\sum_{j=1}^{n} (x_j - x'_j)^2}$$

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Scale first, or vibration decides everything

- Vibration spans **18 Hz** (33 to 51), pressure **2.3 bar** (4.5 to 6.8)
- Unscaled, 10 Hz apart is distance 10 and 1 bar apart is distance 1: vibration picks the neighbours
- The fix (lesson 2): **standardise** each column: subtract its mean, divide by its standard deviation
- In code: `make_pipeline(StandardScaler(), KNeighborsClassifier(5))`

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

k sets the balance between bias and variance

- **Small k**: the vote follows each stored pump, those with flipped labels included. Low bias, high variance (lesson 5)
- **Large k**: the vote reaches pumps across the edge of the envelope, and near the edge the larger class wins. High bias, low variance
- At k = 1 the training accuracy is **exactly 1.000**: each stored pump is its own nearest neighbour

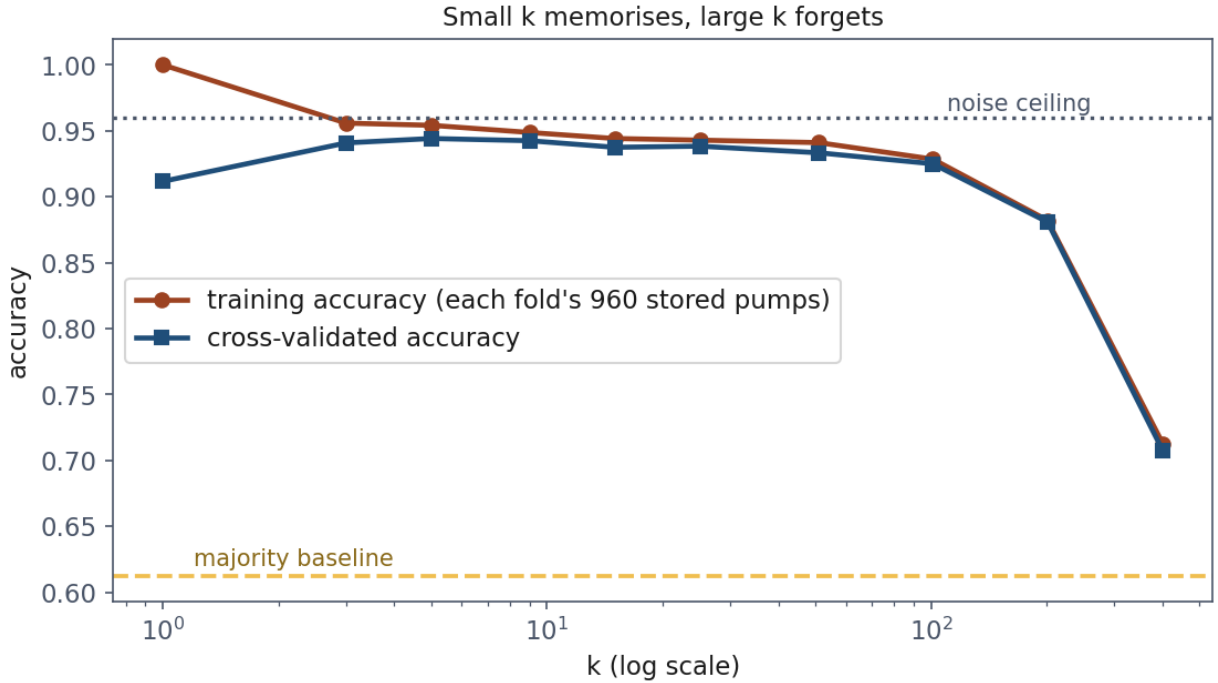
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Training against cross-validated accuracy, by k

k	Training accuracy (on each fold's 960 stored pumps, mean over 5 folds)	Cross-validated accuracy, 5 folds
1	**1.000**	0.912
5	0.954	**0.944**
15	0.944	0.938
51	0.941	0.933
401	0.713	0.708

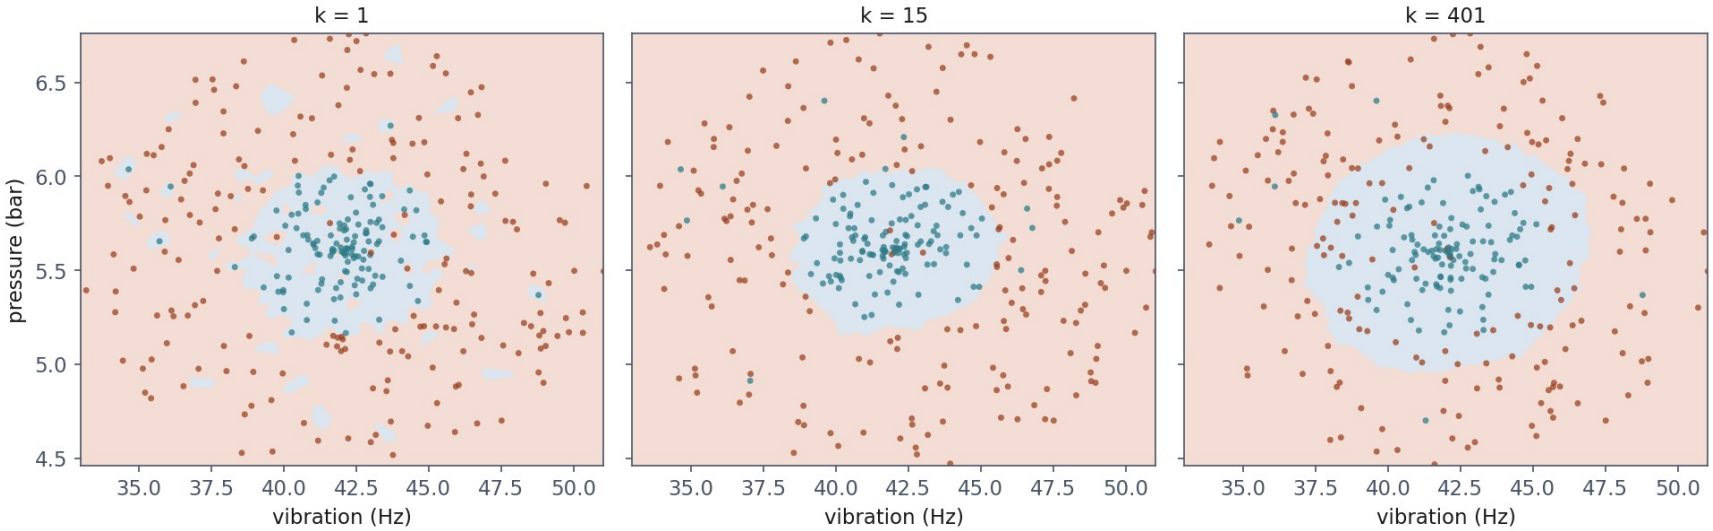
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Accuracy against k, between the baseline and the noise ceiling



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The boundary k-NN draws at k = 1, 15 and 401



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

What k-NN costs

- Storing the pumps costs nothing; the cost comes at **every prediction** instead
- A direct implementation compares each new pump with **every** stored pump: about m × n operations for m pumps and n features
- To predict, the model needs every stored pump: **deploying the model means handing over the data**

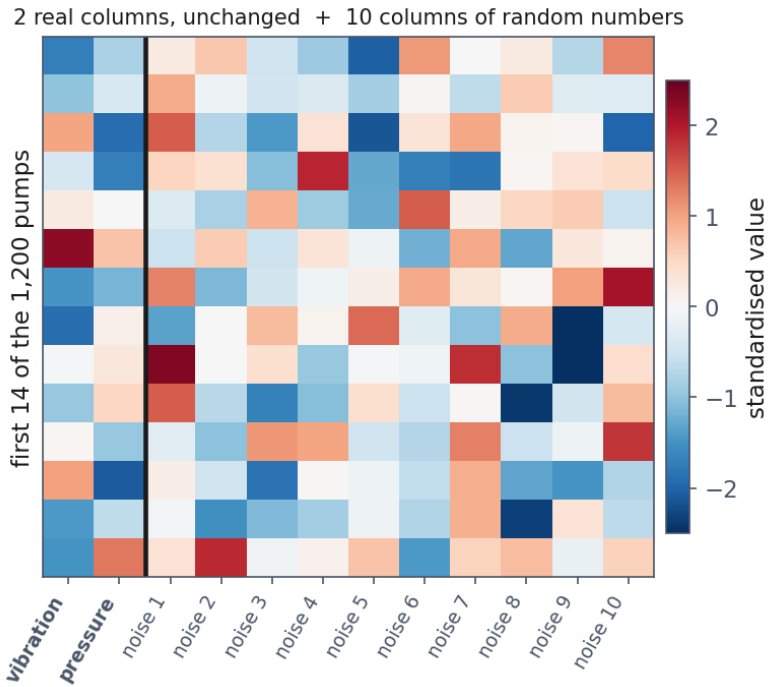
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

An experiment: add columns of pure noise

- Keep the 1,200 pumps and their two real readings **unchanged**
- Add columns of **pure noise**: normal random numbers (mean 0, standard deviation 1), unrelated to the label
- Only the number of columns changes: the problem is as solvable as before
- Same model each time: standardised k-NN, k = 15, 5-fold cross-validation

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The experiment, drawn



The same steps for every number of noise columns

1. Keep the 1,200 pumps and their 2 real readings
2. Add columns of random numbers (mean 0, standard deviation 1), unrelated to the label
3. Standardise, then k-NN with k = 15
4. Score it: 5-fold cross-validated accuracy
5. Repeat with 0, 2, 5, 10, 25, 50, 100 noise columns

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The two real readings are still there, and accuracy still falls

Columns: 2 real + noise	Cross-validated accuracy	Accuracy − baseline 0.613
2	0.938	+0.325
12	0.762	+0.149
27	0.662	+0.049
52	0.602	**−0.011**
102	0.578	−0.035

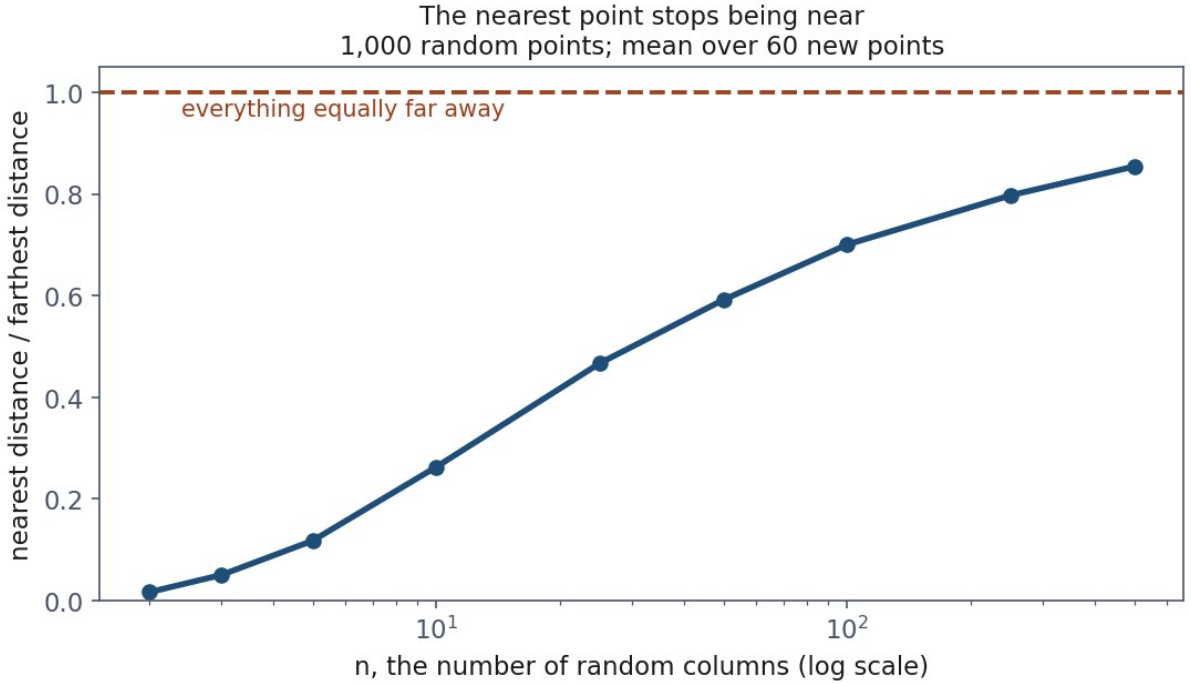
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Why: distances stop varying

- A separate experiment, no pumps: 1,000 random points with n columns, each uniform on [0, 1] (n = 2 a square, n = 3 a cube)
- For a new point, its **nearest** and **farthest** distance: in two dimensions they differ a lot
- Each extra column adds a squared difference to every distance: over many columns **all distances come out about the same**

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

With n random columns, the nearest point stops being near



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Nearest divided by farthest

n, random columns	nearest distance ÷ farthest distance, in n dimensions, mean over 60 new points
2	0.016
10	0.263
50	0.592
100	**0.701**
500	0.855

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

In 100 dimensions the nearest point is 70% as far as the farthest

- In **two** dimensions the ratio is 0.016: the nearest point is at **2%** of the farthest distance
- In **one hundred** it is 0.701, **70%**: the nearest point is barely nearer than a random one
- Back to the pumps: the 15 nearest are almost 15 pumps at random, so with 52 columns accuracy fell below the baseline

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

What the curse is, and is not

- **Not** that many columns cannot be learned: a linear model ignores a useless one
- It is that **distance-based methods stop working**: all distances come out alike
- For k-NN: choose the columns first, inside the pipeline
- In code: `SelectKBest(mutual_info_classif, k=2)` before k-NN: back to **0.938**

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Notebook 1, live

section, and its slides	you call (→ the number to come back with)
the straight line, 4-6	`LogisticRegression`, `StratifiedKFold`, `cross_val_score` → 0.613, `cross_val_predict`
k-NN, 7-10	by hand = `KNeighborsClassifier`; then `make_pipeline`, `StandardScaler`
choosing k, 11-14	`cross_validate` with training scores → 0.944 at k = 5
columns of noise, 16-18	k = 15 on `load_with_noise(50)` → 0.602
distances, 19-22	NumPy: nearest ÷ farthest → 0.701 at n = 100
choosing the columns, 23	`SelectKBest`: `f_classif` 0.562, `mutual_info_classif` 0.938
or atoms 1-5	`Consolidation/06_lesson_6_atoms.ipynb`

Break

- Twelve minutes

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Naive Bayes: ask how likely the readings are in each class

- A second family: we want P(y = c | x), the probability of class c **given** the readings x
- Bayes' rule uses how common each class is, and which readings are common in each (healthy near 42 Hz and 5.6 bar, faulty all around)
- The hard part: how often a class shows **this exact combination** of readings - with many readings, almost no pump of that class is near x

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Bayes' rule, applied to a class

$$P(y = c \mid x) = \frac{p(x \mid y = c)\ P(y = c)}{p(x)}$$

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Drop the denominator: arg max picks the class with the largest numerator

$$\hat{y} = \arg\max_c \, p(x \mid y = c) \, P(y = c)$$

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The assumption: independent, given the class

- To make p(x | y = c) estimable, Naive Bayes assumes: **given the class, the features are independent of one another**
- Then the density of all n readings together is the product of n one-dimensional densities, one per feature
- Each one-dimensional density is estimated from the pumps of that class alone. In code:

```
GaussianNB()
```


UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The assumption, written: a product of n densities

$$p(x \mid y = c) = \prod_{j=1}^{n} p(x_j \mid y = c)$$

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Gaussian: each factor is a bell curve

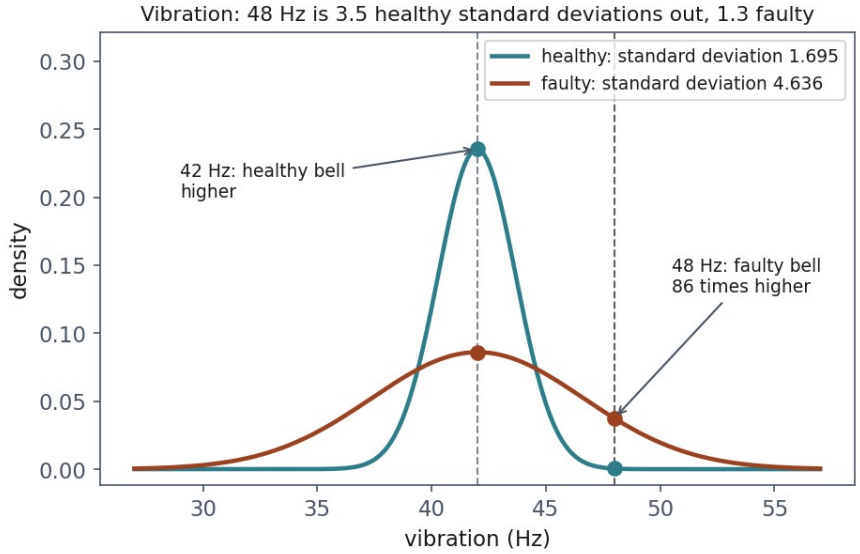
$$p(x_j \mid y = c) = \frac{1}{\sqrt{2\pi\sigma_{jc}^2}} \exp\left[-\frac{(x_j - \mu_{jc})^2}{2\sigma_{jc}^2}\right]$$

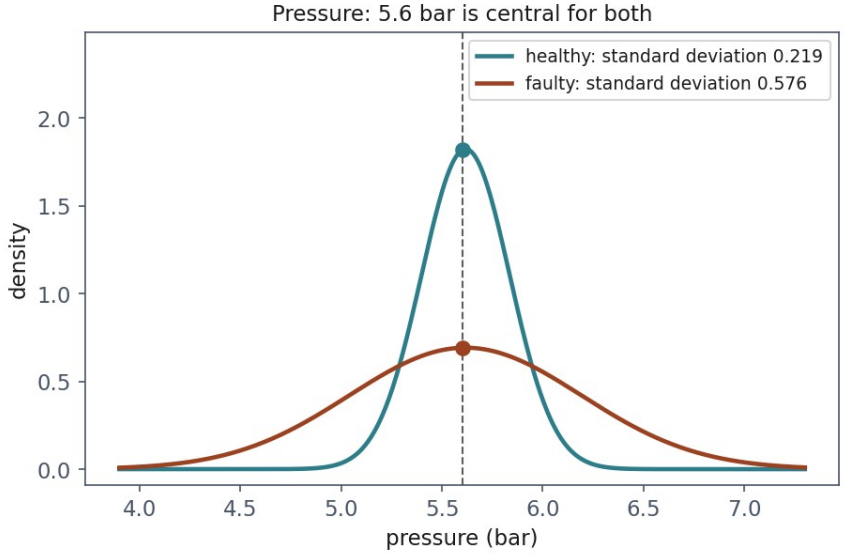
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Training computes ten numbers: the widths separate the classes

Fitted on the 1,200 pumps	Healthy	Faulty
Prior: the class's share of the pumps	465 ÷ 1,200 = 0.3875	735 ÷ 1,200 = 0.6125
Vibration mean	41.986 Hz	42.015 Hz
Vibration standard deviation	**1.695 Hz**	**4.636 Hz**
Pressure mean	5.617 bar	5.616 bar
Pressure standard deviation	**0.219 bar**	**0.576 bar**

A pump at 48 Hz, 5.6 bar: there the faulty bell is 86 times higher





UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Logs turn the product into a sum, and keep the same winner

$$\ln\left[P(y=c)\prod_j p(x_j \mid y=c)\right]$$

$$= \ln P(y=c) + \sum_j \ln p(x_j \mid y=c)$$

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

48 Hz, 5.6 bar: scores and their logs

Class	Row	Prior	Vibration density	Pressure density	Product of the three, or sum of their logs
Healthy	value	0.3875	0.000433	1.818	0.000305
Healthy	ln	−0.948	−7.745	+0.598	−8.095
Faulty	value	0.6125	0.0374	0.692	0.01585
Faulty	ln	−0.490	−3.286	−0.368	**−4.144**

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

From the two totals to probabilities

Step	Healthy	Faulty
1. Log total, previous slide	−8.095	−4.144
2. Undo the log: e^total = prior × densities	0.000305	0.01585
3. Divide each by their sum, 0.000305 + 0.01585 = 0.01616 = p(x)	**0.019**	**0.981**

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Software uses only the gap between the logs

- Many features: e^total rounds to 0 for both classes, and 0 ÷ 0
- The gap stays ordinary: −4.144 − (−8.095) = 3.951
- Divide top and bottom by e^(faulty total): P(faulty) = 1 / (1 + e^(healthy total − faulty total))
- P(faulty) = 1 / (1 + e^(−3.951)) = **0.981**: odds of e^3.951 ≈ **52 to 1**

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

A cheap model, built on an assumption that is almost never true

- The gain: training is **one pass** over the data; extra features cost almost nothing
- The assumption rarely holds: on a real pump, vibration and pressure both depend on how hard it runs
- In a text, “New” and “York” appear together far more than independence allows
- So the useful question: **when does being wrong about it cost nothing?**

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Naive Bayes on the pumps: 0.933

Model	Accuracy, 5-fold cross-validation, same folds	Standard deviation over the 5 folds
Majority baseline	0.613	0.000
Logistic regression	0.613	0.000
Gaussian Naive Bayes	**0.933**	0.010
k-NN, k = 5	0.944	0.006

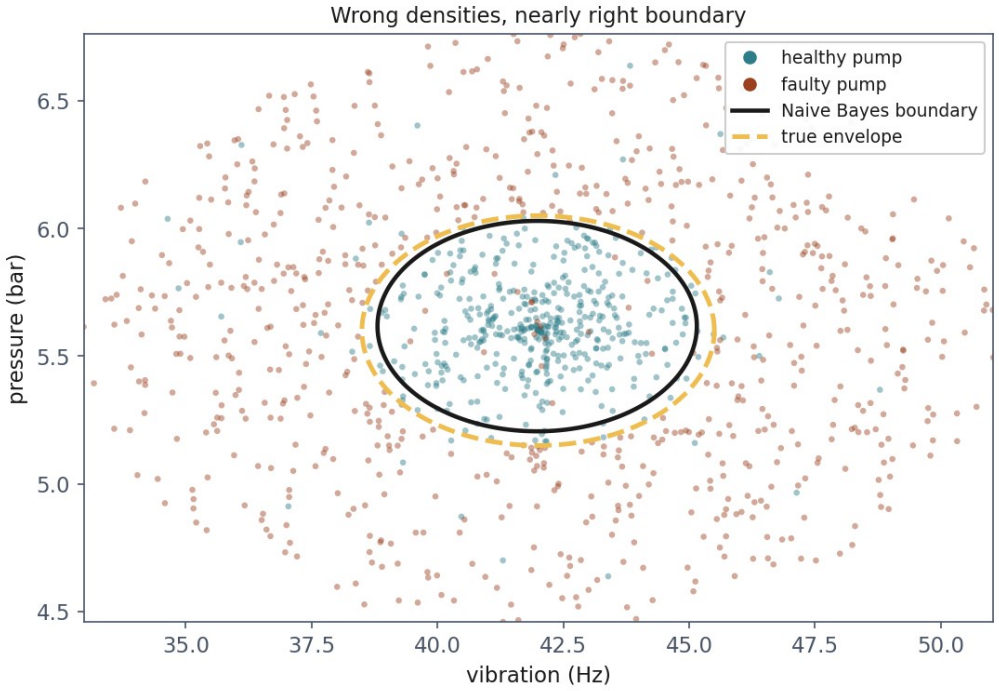
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Uncorrelated within each class, but not independent

Correlation, within one class	between vibration and pressure	between the two standardised readings, squared (how far from the centre)
healthy pumps	−0.006	+0.168
faulty pumps	−0.049	**−0.420**

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The densities are wrong, but the boundary they draw is about right



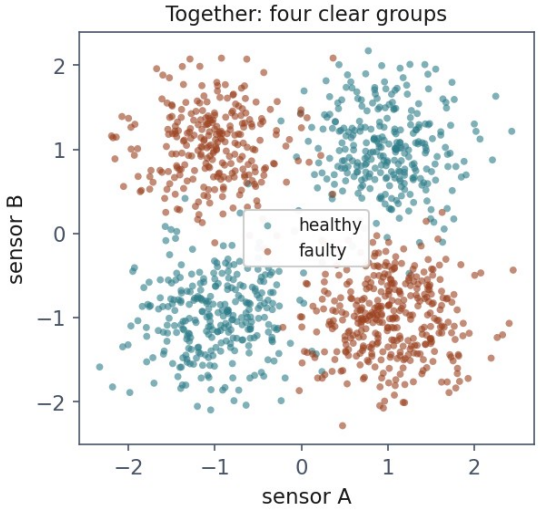
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

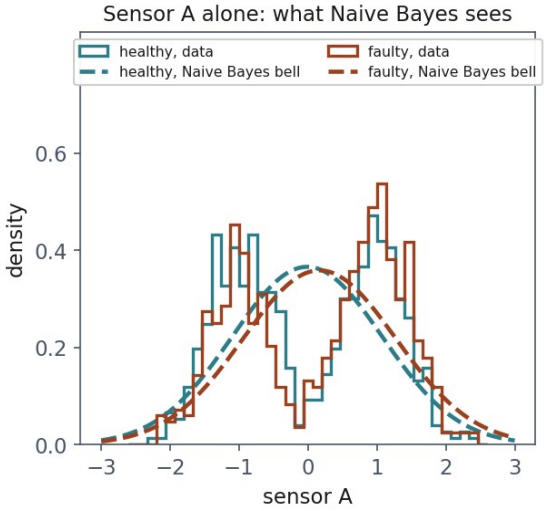
A second dataset: the signal is an interaction

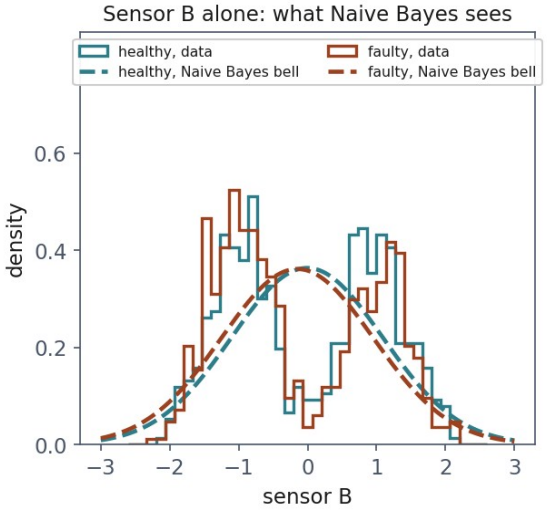
- 1,200 more simulated pumps, with two other sensors, A and B
- Each sensor is either high (around +1) or low (around −1), half the pumps each; faulty when **exactly one** is high - the exclusive or (XOR)
- Both high: high-load mode. Both low: idle. One alone: demand and delivery do not match
- So A alone says nothing: what a high A means depends on B (an **interaction**)

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Each sensor alone: the two classes look the same







UNIVERSITÀ DEGLI STUDI DELL'AQUILA

0.404: below the baseline, and below chance

Model	Accuracy, 5-fold cross-validation
Majority baseline: 628 faulty ÷ 1,200	0.523
Gaussian Naive Bayes	**0.404**
Logistic regression	0.393
SVM, linear kernel	0.606
k-NN, k = 5	0.967
SVM, radial basis function (RBF) kernel	0.972

Its probabilities cannot be trusted

- Interacting sensors: Naive Bayes fitted on 900 pumps, tested on the other 300
- For each pump it gives P(faulty) and predicts the larger class. Confidence = the probability of the class it predicts, from 0.5 (a coin) to 1 (certain)
- Mean confidence **0.567** when right, **0.555** when wrong: it cannot tell when it is wrong
- It says about 56% sure and is right on 133 of 300, 44%: its **probabilities are not calibrated**

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Notebook 2, live

section, and its slides	you call (→ the number to come back with)
on the pumps, 39	`GaussianNB` beside `KNeighborsClassifier` and `LogisticRegression` → 0.933
what training computed, 31-37	`.theta_`, `.var_`, `.class_prior_`, `predict_proba` → 0.981 at 48 Hz
checking the assumption, 40-41	squared standardised readings, `.corr()` within each class → −0.420
the interacting sensors, 42-44	the same models, plus `SVC` → 0.404
confidence, right and wrong, 45	`train_test_split`, then `predict_proba` → 0.567 against 0.555
or atoms 6-10	`Consolidation/06_lesson_6_atoms.ipynb`

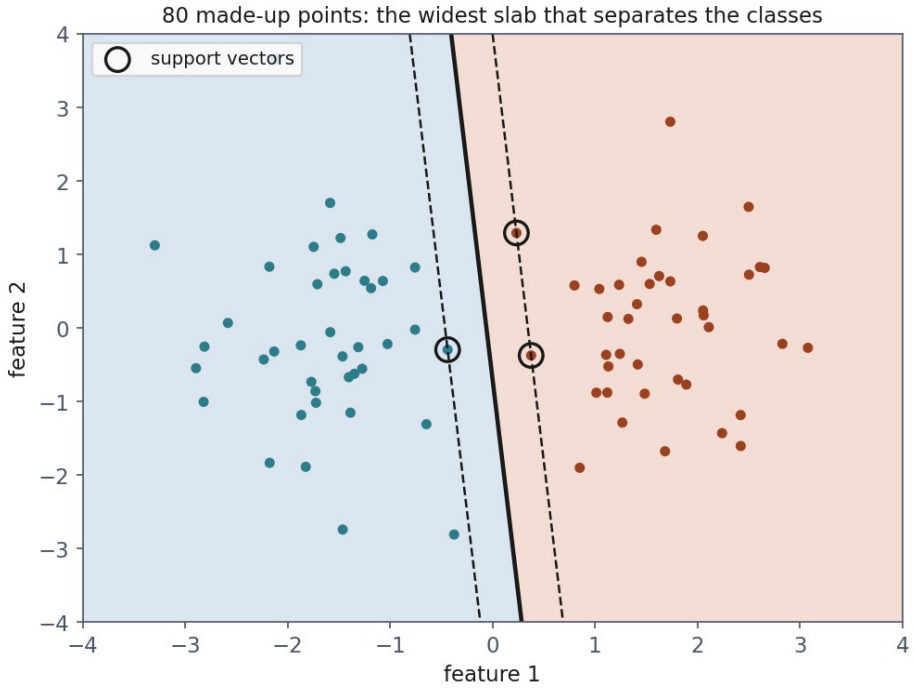
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

SVM: among the lines that fit, take the widest margin

- When a line can separate two classes, **infinitely many** lines do, all with **zero training error**: the error cannot choose
- The SVM picks the line with the widest **margin**: an empty strip (the slab) around the line, widened until it touches a point of each class
- The reason: a line that passes close to a point gets it wrong after a small change in that point's readings

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

80 made-up points: three of them fix the line (support vectors)



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The strict problem: the widest slab with no point inside

- Labels -1 (healthy), $+1$ (faulty): $y \cdot f(x) > 0$ means “classified right”; in code y stays 0/1
- Boundary $w \cdot x + b = 0$, scaled so the closest points have $|w \cdot x + b| = 1$: the slab is $2/\|w\|$ wide
- Widest slab = shortest w: minimise $\frac{1}{2}\|w\|^2$, every point outside, $y_i(w \cdot x_i + b) \geq 1$
- No term for errors: the constraints forbid them

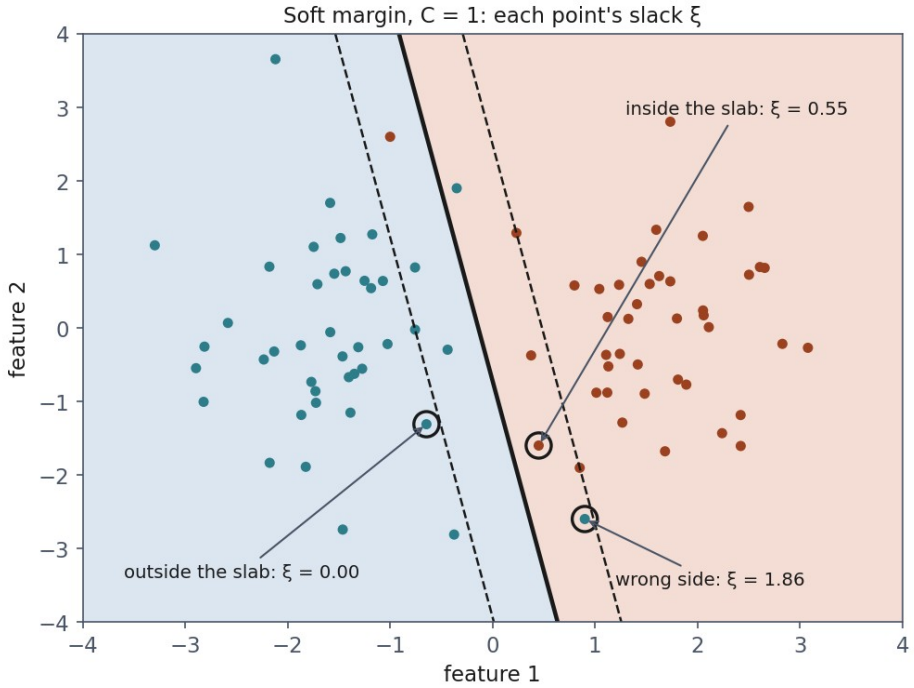
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Not separable? Charge for the violations

- With 4% of the labels flipped, no slab separates the pumps: the strict problem has **no solution**
- Give each point i an allowance $\xi_i \geq 0$, its slack: how far it may fall short of its edge of the slab. Each unit of slack costs C
- $\xi_i = 0$ outside the slab, 0 to 1 inside it, above 1 on the wrong side
- That charge is the **hinge loss** $\max(0, 1 - y \cdot f(x))$, labels ± 1, $f(x) = w \cdot x + b$

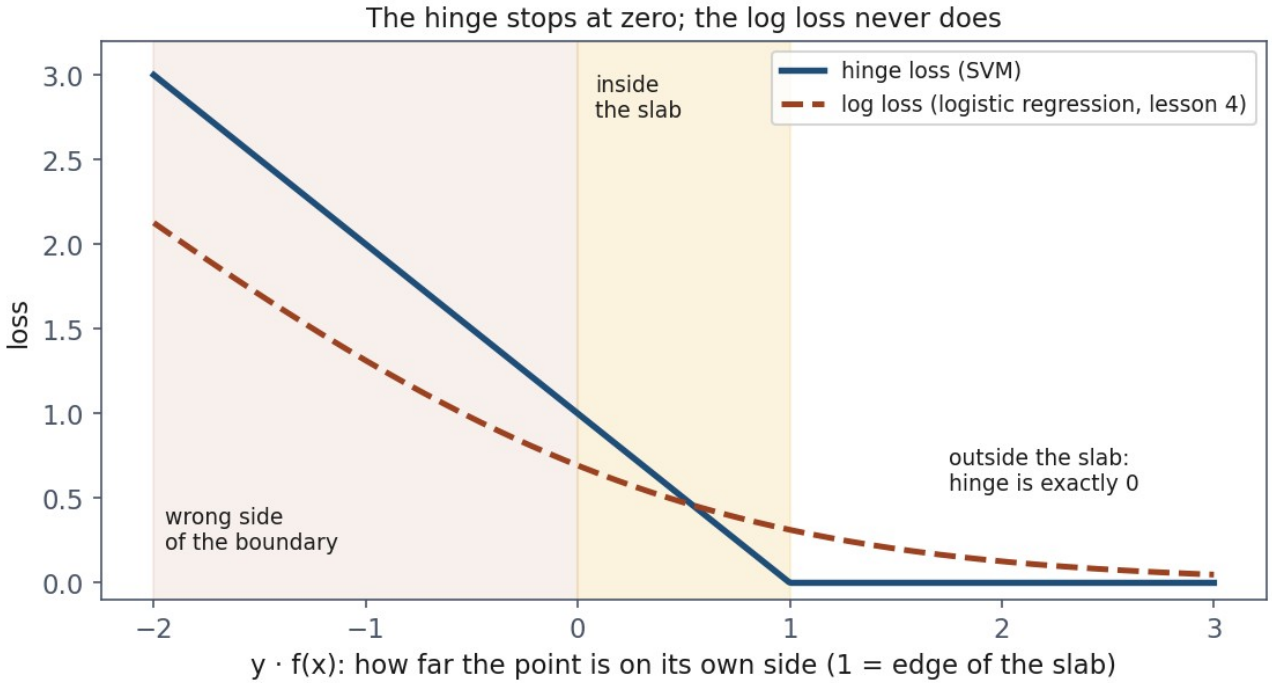
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Three kinds of point: outside the slab, inside it, on the wrong side



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The hinge loss stops at zero



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The SVM objective: small w, plus C times the slack

$$\min_{w,b,\xi} \frac{1}{2}\|w\|^2 + C\sum_i \xi_i$$

subject to $y_i(w^\top x_i + b) \geq 1 - \xi_i,\ \xi_i \geq 0$

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

C is the price of a training error

- **Large C**: slack is expensive, the slab narrows to classify more training points. Low bias, high variance
- **Small C**: slack is cheap, the slab stays wide, some training points stay wrong. Higher bias, lower variance
- C plays the role of k in k-NN and of λ in lesson 3, but **large C means less regularisation**: divided by C, the penalty weighs 1/(2C)

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

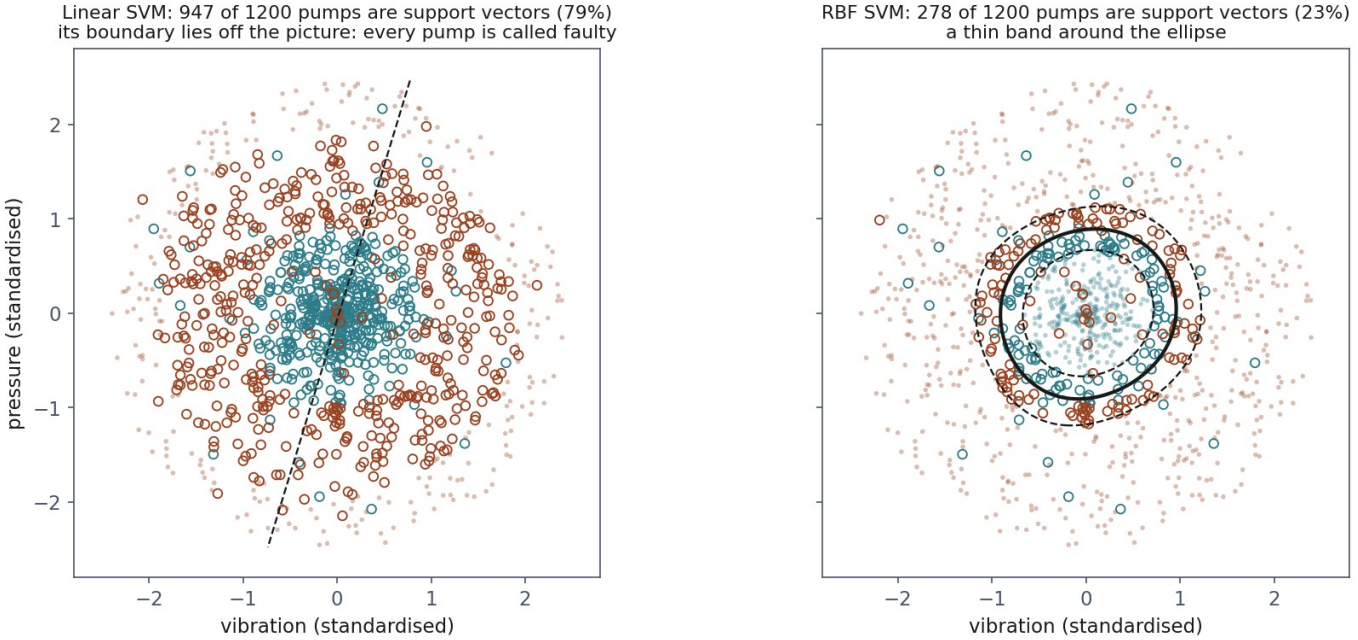
The best straight line is still a straight line

Model, on the pumps (standardised, default C = 1)	Accuracy, 5-fold cross-validation	Support vectors, fitted on all 1,200
SVM, linear kernel: a straight line	**0.613 ± 0.000**	947 of 1,200 (**79%**)
SVM, RBF kernel: a curved boundary	0.947 ± 0.005	278 of 1,200 (23%)

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

On the pumps: a straight line leaves almost every pump in its slab

Circled: support vectors, on the edge of the slab, inside it, or on the wrong side. Solid: boundary; dashed: edges of the slab



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Many support vectors: a free warning

- The linear SVM, `SVC(kernel="linear")`, keeps **79%** of the pumps as support vectors; the RBF SVM keeps **23%**
- 79% means almost every pump sits on or inside the slab: no strip of empty space separates the classes
- The count comes free after `fit`, as `len(model.support_)`: a high share is a warning before any cross-validation

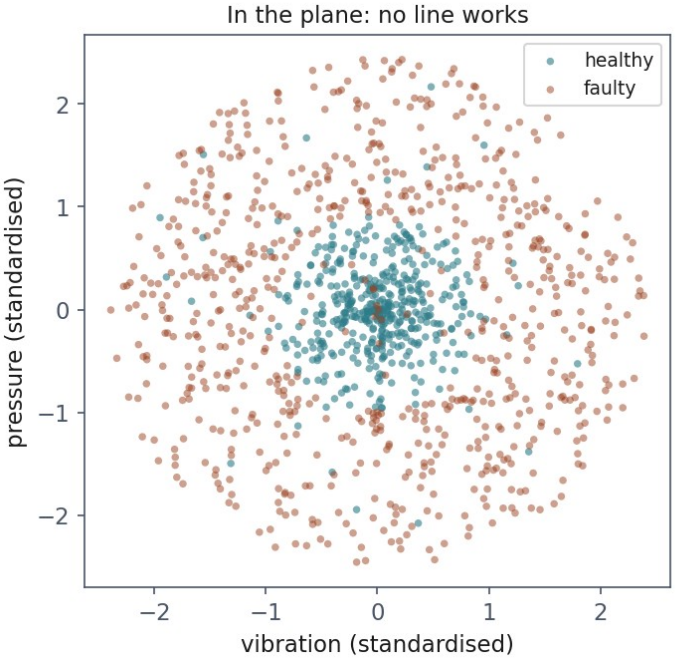
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

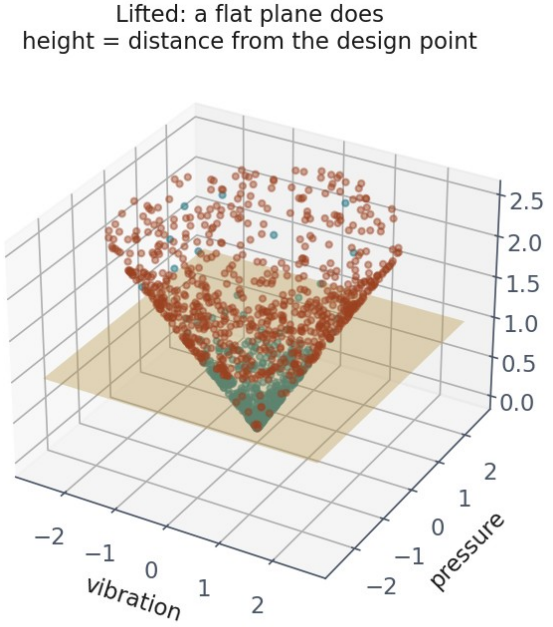
A new coordinate makes the classes separable

- In the plane, healthy pumps sit in a disc, faulty ones around it: no line separates them
- Add a third coordinate: each pump’s **distance from the design point**, standardised
- Healthy pumps rise a little, faulty ones a lot: a flat plane separates them, flipped labels apart
- The rule was separable all along, in other coordinates

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The same pumps, lifted





UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The SVM’s solution: a weighted sum of pumps

- Slide 53’s objective: $\frac{1}{2}\|w\|^2 + C \Sigma \max(0, 1 - y_i(w \cdot x_i + b))$
- Its slope in w: $w - C \Sigma y_i x_i$, over the pumps inside the slab or on the wrong side
- Zero slope: $w = \Sigma a_i y_i x_i$; $a_i = C$ inside, 0 outside, 0 to C on the edge
- Only the support vectors have $a_i > 0$: the pegs of the tent

UNIVERSITÀ DEGLI STUDI DELL’AQUILA

So the new coordinates are never computed

- $\varphi(x)$: a pump's new coordinates. Slide 60 there: $w = \Sigma\, a_i\, y_i\, \varphi(x_i)$
- Boundary $\Sigma\, a_i\, y_i\, \varphi(x_i)\cdot\varphi(x) + b$; training needs only $\varphi(x_i)\cdot\varphi(x_j)$: the pumps enter **only through inner products**
- A **kernel** $K(x, x')$ gives $\varphi(x)\cdot\varphi(x')$ from the original readings: φ is never computed
- RBF, scikit-learn's default: one exponential per pair

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The kernel trick by hand: 25 = 25

Route, for u = (1, 2) and v = (3, 1)	Computation	Result
Lift with $\varphi(x) = (x_1^2, \sqrt{2}\cdot x_1 x_2, x_2^2)$, then multiply	$\varphi(u) \cdot \varphi(v) = (1, 2.83, 4) \cdot (9, 4.24, 1) = 9 + 12 + 4$	**25**
Never lift: the kernel $(u \cdot v)^2$	$(1\times3 + 2\times1)^2 = 5^2$	**25**

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The RBF kernel: similarity falls with distance

- Radial basis function (RBF) kernel: $K(x, x') = \exp(-\gamma\|x - x'\|^2)$
- A similarity: 1 for identical pumps, towards 0 far apart. With $\gamma = 0.5$: 0.61 at one standard deviation, 0.14 at two
- Why it: expand the exponential and it holds $(x \cdot x')^2$, $(x \cdot x')^3$, … every degree at once - infinitely many new coordinates, no lift to guess
- Cost: one exponential per pair of pumps

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Training: one weight per pump

- fit minimises slide 53’s objective in the weights a_i, two at a time: no learning rate, one answer (RBF, $\gamma = 0.5$, $C = 1$)
- **922** pumps get weight 0: outside the slab, they drop out
- **268** get $C = 1$ (inside the slab or wrong side), **10** sit on its edge: 278 support vectors
- The model kept: those 278 pumps, their weights, and $b = +1.105$

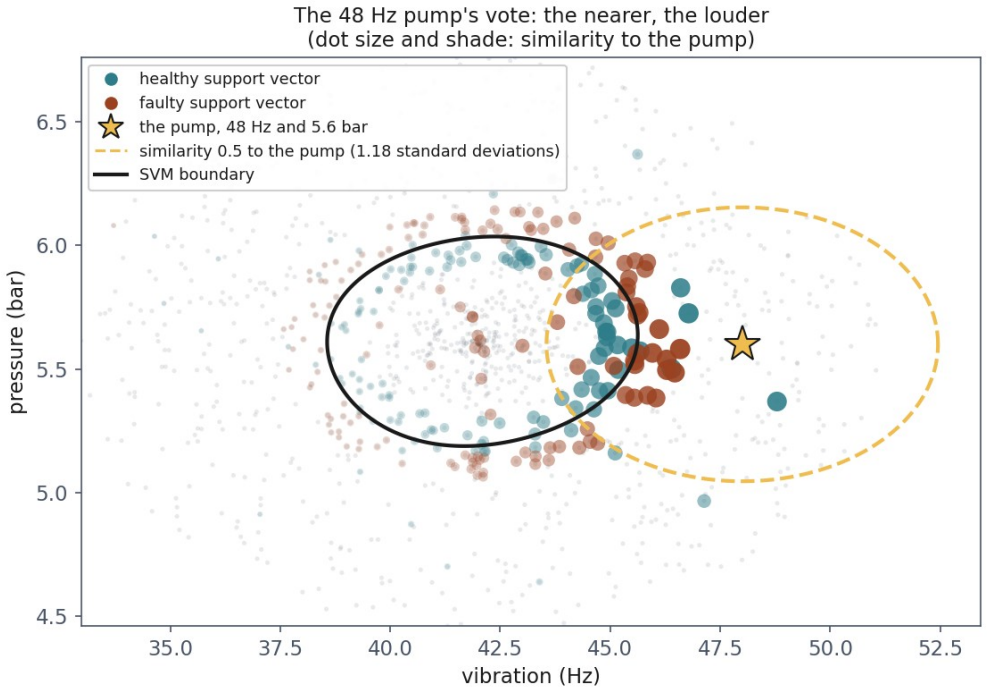
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The SVM’s score at 48 Hz, 5.6 bar: $\Sigma\, a_i\, y_i\, K(x_i, x) + b = +1.883$

Term, for the pump x at 48 Hz, 5.6 bar	Value
Faulty support vectors: sum of $a_i \cdot K(x_i, x)$	+37.455
Healthy support vectors: minus the sum of $a_i \cdot K(x_i, x)$	−36.677
Intercept b	+1.105
Total: positive, so faulty	**+1.883**

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

One pump’s vote, drawn



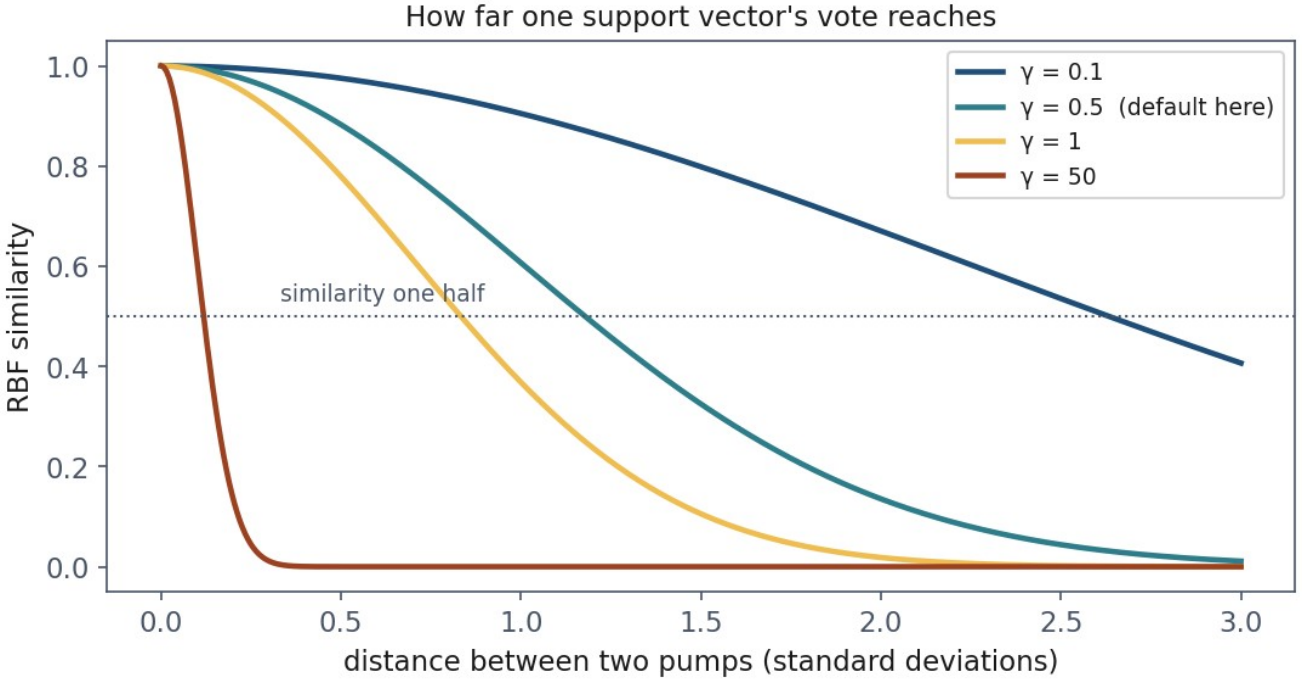
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Training against cross-validated, three settings

RBF SVM on the pumps	Training accuracy (each fold's 960 pumps)	Cross-validated accuracy, 5 folds
$\gamma = 0.1, C = 1$	0.934	0.929
$\gamma = 1, C = 1$	0.948	**0.944**
$\gamma = 50, C = 1000$	**0.996**	0.902

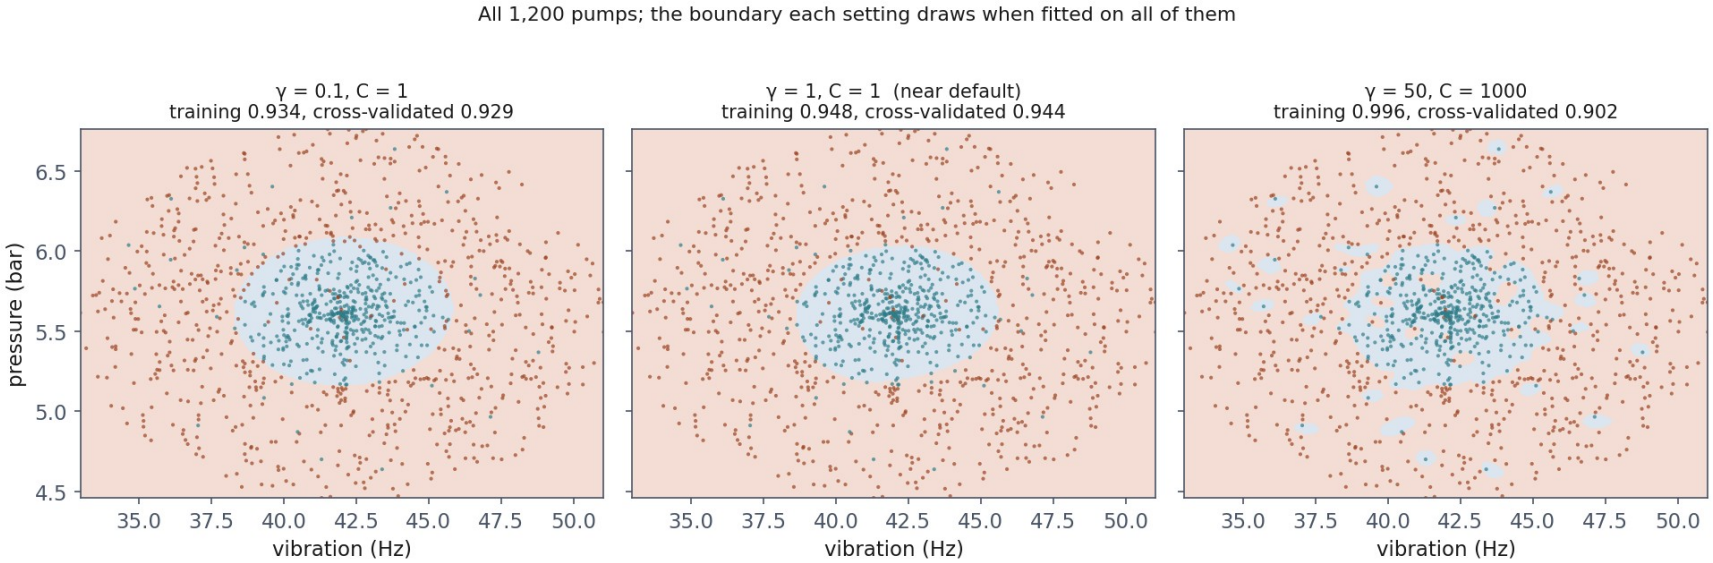
UNIVERSITÀ DEGLI STUDI DELL'AQUILA

How far one support vector’s vote reaches, for four values of γ



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Overfitting you can see: the boundary at the three settings



UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Notebook 3, live

section, and its slides	you call (→ the number to come back with)
the widest slab, 47-52	`SVC(kernel="linear")`, `.n_support_`, `decision_function` → 3 of 80 points
a straight line still fails, 55-57	`cross_val_score`, linear against RBF kernel → 947 against 278 support vectors
the kernel trick, 58-63	NumPy only: $(u \cdot v)^2$ against the lifted product → 25 = 25
training and one prediction, 64-66	`.support_`, `.dual_coef_`, `.intercept_` → +1.883
γ and C, 67-69	`SVC(gamma=..., C=...)`, `cross_validate` → 0.944 at γ = 1
or atoms 11-15	`Consolidation/06_lesson_6_atoms.ipynb`

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

The three families compared, on the same pumps and the same folds

Model	Accuracy, 5-fold cross-validation	Standard deviation over the 5 folds
Majority baseline	0.613	0.000
Logistic regression (lesson 4)	0.613	0.000
SVM, linear kernel	0.613	0.000
Gaussian Naive Bayes	0.933	0.010
k-NN, k = 5	0.944	0.006
SVM, RBF kernel	**0.947**	0.005
noise ceiling	*≈ 0.96*	

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

How to choose, in practice

Situation	Reach for
Few features, plenty of data, a curved boundary	k-NN, or an RBF SVM
Very many features, little data	Naive Bayes
A line separates the classes	Logistic regression or a linear SVM
You need probabilities you can trust	Check calibration on held-out data; logistic regression is the usual start; **not** Naive Bayes; an SVM gives only a score
The model must be small, or fast at prediction	Naive Bayes; an SVM with few support vectors; not k-NN
The signal is an interaction between features	k-NN or an RBF SVM, not Naive Bayes

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

What to take away

- **No line separates the pumps**: 0.613
- **k-NN**: k sets bias against variance; in 100 random columns the nearest is 70% as far as the farthest
- **Naive Bayes**: 0.933, but **0.404** on interactions
- **SVM**: support vectors decide; kernels bend the line
- **Choosing the family beats tuning**: 0.613 to 0.947

UNIVERSITÀ DEGLI STUDI DELL'AQUILA

Homework

- **Exercise 6**, discussed at the start of **Friday 27 November**
- Two stations: fit five models from the three families and the linear baseline, and argue which one you would choose, and why